

COMP20081 SYSTEMS SOFTWARE

REVISION GUIDE

BY MATT ROBINSON

Copyright © 2017 Revision Guide

Copying prohibited

All rights reserved. No part of this publication may be reproduced or transmitted in any form or by any means, electronic or mechanical, including photocopying and recording, or by any information storage or retrieval system, without the prior written permission of the publisher.

Art. No xxxxx

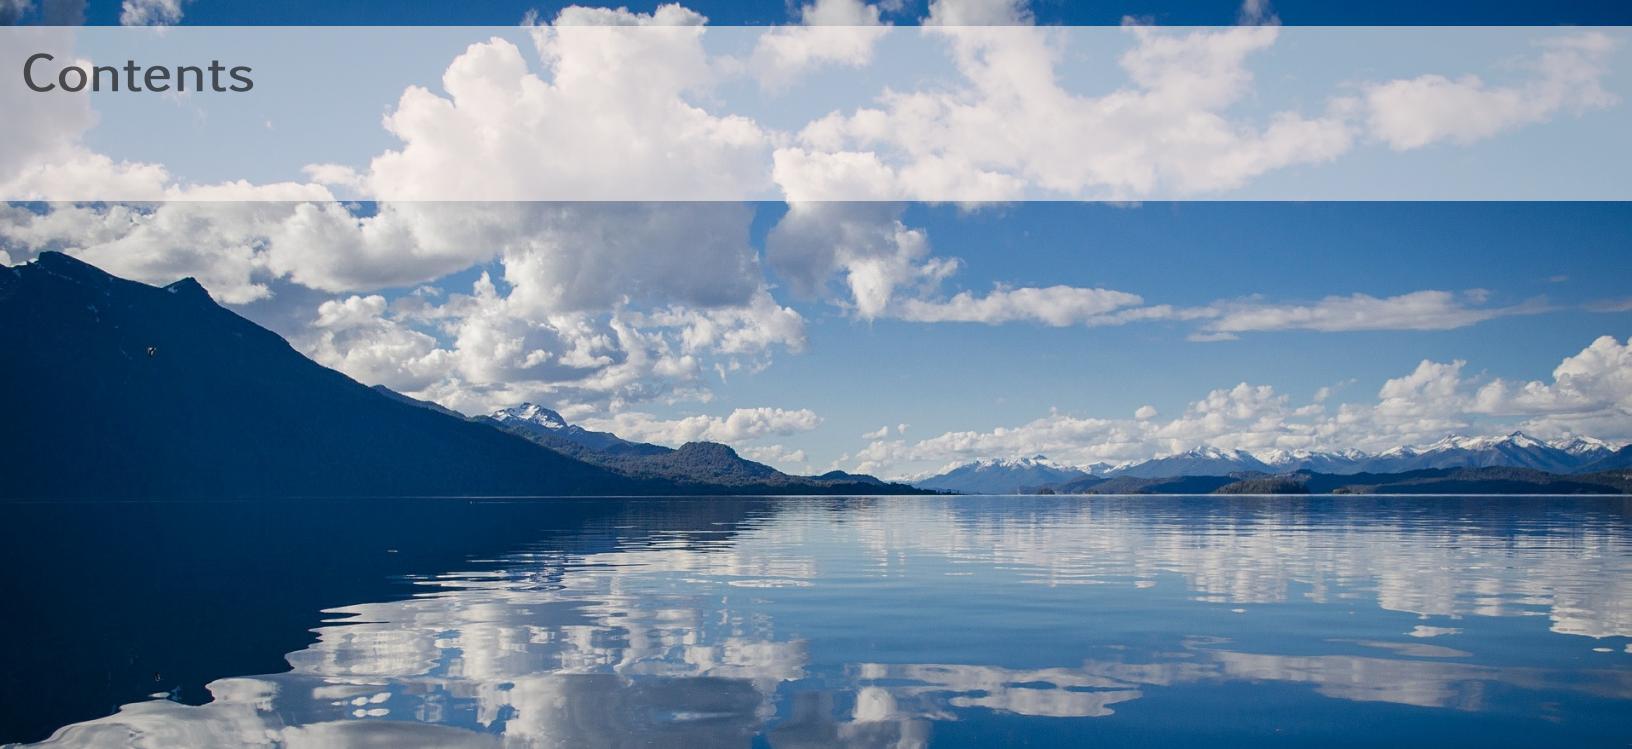
ISBN xxx-xx-xxxx-xx-x

Edition 0.0

Cover design by Cover Designer

Published by by Matt Robinson
Printed in City

Contents



| | | |
|----------|---------------------------------------|-----------|
| 1 | Introduction | 5 |
| 1.1 | History of performing computations | 5 |
| 1.2 | Hardware | 7 |
| 1.3 | What is an operating system (OS)? | 9 |
| 1.4 | Functions of an operating system | 11 |
| 1.5 | Current operating system (OS) trends | 13 |
| 1.6 | Operating system (OS) layers | 15 |
| 1.7 | Kernel mode | 16 |
| 2 | Process Management | 19 |
| 2.1 | Programs and processes | 19 |
| 2.2 | Process life cycle | 20 |
| 2.3 | Program execution | 26 |
| 2.4 | Concurrency | 28 |
| 2.5 | Process scheduling | 29 |
| 3 | Threads and Concurrency | 39 |
| 3.1 | What are threads? | 39 |
| 3.2 | Sequential and concurrent programming | 43 |
| 3.3 | Sequential execution | 45 |
| 3.4 | Introduction to concurrent execution | 47 |

1. Introduction

| | | |
|-----|--------------------------------------|----|
| 1.1 | History of performing computations | 5 |
| 1.2 | Hardware | 7 |
| 1.3 | What is an operating system (OS)? | 9 |
| 1.4 | Functions of an operating system | 11 |
| 1.5 | Current operating system (OS) trends | 13 |
| 1.6 | Operating system (OS) layers | 15 |
| 1.7 | Kernel mode | 16 |

1.1 History of performing computations

1950s

A user may want to perform computations that may have a complexity that is not feasible by a standard calculator.

A description of the computations to be performed could be expressed in a physical form using a perforated card.

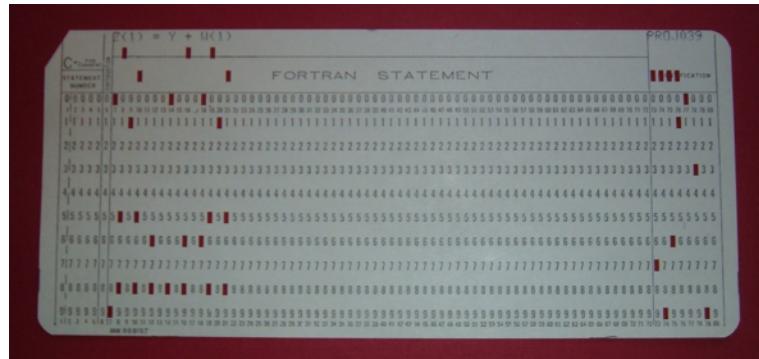


Figure 1.1: Perforated card

There are multiple lines on the card, each with varying holes. Each line describes a particular instruction that is necessary for the computations to be completed.



>1960s

The process of performing computations after 1960 changed dramatically and required far less human input.

Automation was brought about by the “Operating System Paradigm” in which users could interface with a computer system directly through an operating system (OS).



History details

| | |
|------------|--|
| Late 1950s | <ul style="list-style-type: none"> • Standard subroutines were produced that were loaded at start-up. These contained features similar to those found on an operating system. • Magnetic tapes were used for storage and were later replaced by disks. • Assemblers started to be used. These are programs that takes basic computer instructions and convert them in to machine code; this is a pattern of binary bits (0's and 1's) that the computer system's processor can use to perform its basic operations. • High-level languages, which consisted of more natural and human-readable language, started to be used. For example, FORTRAN is a general-purpose, compiled imperative programming language that is especially suited to numeric computation and scientific computing and was introduced in 1957. |
|------------|--|

| | |
|-------|--|
| 1960s | Automated batch system. <ul style="list-style-type: none">• This replaced the computer operator.• Several programs could be loaded into memory and automatically processed in a “first in, first out” (FIFO) fashion. |
| 1970s | Multiprogramming. <ul style="list-style-type: none">• The computer could switch between jobs, which allows processing and input/output (I/O) interaction simultaneously. |
| 1980s | Graphical user interfaces (GUIs). <ul style="list-style-type: none">• The interaction between a computer system and a user through the medium of a mouse and keyboard. |

1.2 Hardware

External hardware

A **peripheral** is any external hardware device that provides input/output (I/O) for the computer.

For example, a keyboard and mouse are input peripherals, while a monitor and printer are output peripherals. Some peripherals, such as external hard drives, provide both input and output for the computer.

A computer system generally has many internal hardware components and hardware peripherals.



Figure 1.4: External view of computer system

Internal hardware

A **processor** or **central processing unit (CPU)** is the hardware within a computer that carries out the instructions of a computer program by performing the basic arithmetical, logical, and input/output operations of the system.

A **motherboard** is the main printed circuit board (PCB) in a computer. The motherboard is a computer's central communications backbone connectivity point, through which all components and external peripherals connect.

Inside of a computer system, there are many components connected to the processor via the motherboard.

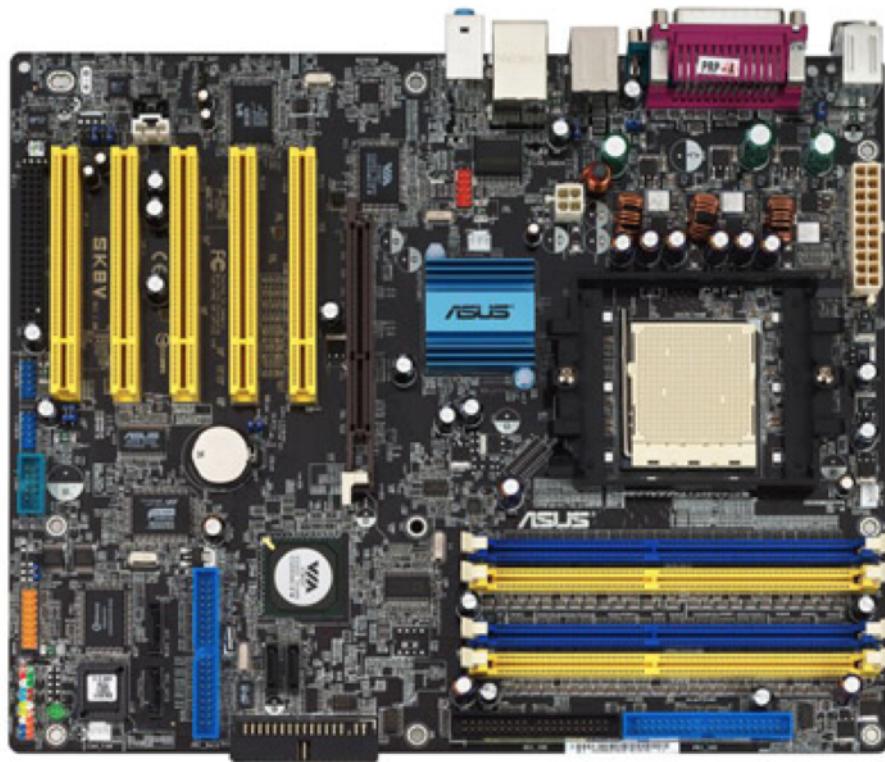


Figure 1.5: Typical computer motherboard

The motherboard connects:

- all of the internal components via the data bus; and
- the peripherals.

1.3 What is an operating system (OS)?

Definition

An **operating system (OS)** is a collection of system programs that manage the hardware resources and peripherals connected to a computer system. It is also responsible for the graphical user interface or command line interface and all other software running on the computer system.

Purpose of an operating system (OS)

An operating system (OS) is designed to:

- eliminate the need to have hardware knowledge to operate a computer system;

- make the boundary between hardware and software transparent, allowing the user to not be concerned with the technical details; and
- provide a user-friendly environment to execute and develop programs.

These attributes are achieved by layering the computer system such that the user can interface with applications, rather than the operating system (OS) or the hardware directly.

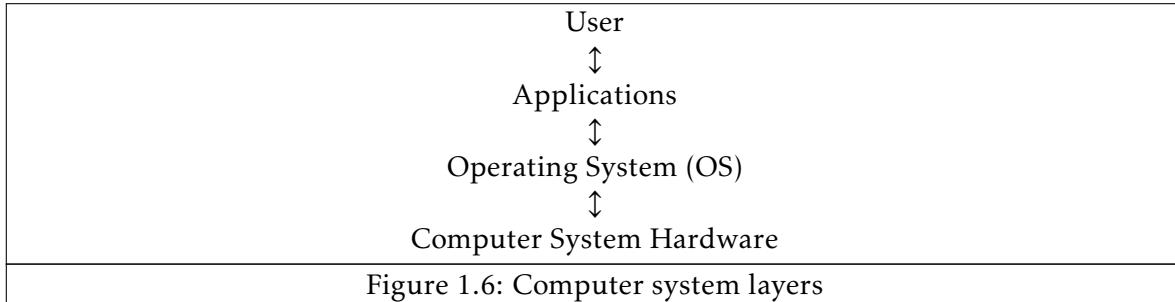


Figure 1.6: Computer system layers

Structure of an operating system (OS)

The structure of an operating system (OS) can be said to resemble an onion.

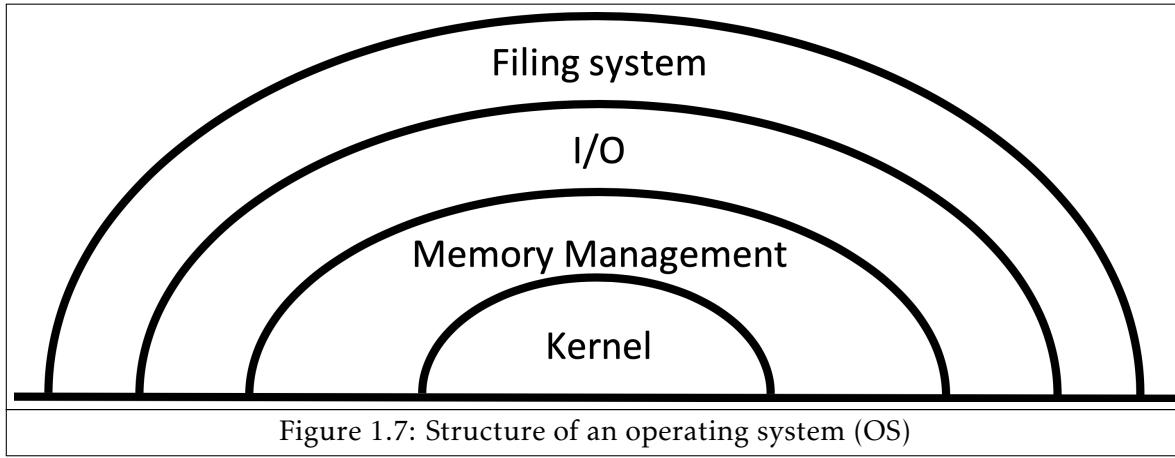


Figure 1.7: Structure of an operating system (OS)

An operating system (OS) has four main components.

The **kernel** hides the complexity of how a computer system works from users. It is responsible for:

- process management;
- CPU scheduling; and
- handling interrupts.

Memory management is responsible for allocating and deallocating memory to processes.

Input/output (I/O) includes any interaction between the internal computer system components and peripherals.

The **filing system** is comprised of file management subsystems.

Each layer in the operating system (OS) structure provides functions to the above layers.
Each layer uses facilities provided by layers within and below that layer.

Practical features of current operating systems (OSs)

| | |
|-----------------------------|--|
| Concurrency | Allows overlapping input/output (I/O) operations with computations and several programs to be stored in memory at a single time. |
| Sharing of resources | Sharing hardware and peripherals, such as hard disks and printers. |
| Access to long term storage | Important for saving important files on mediums such as hard disk drives (HDDs) and solid-state drives (SSDs). |
| Non-determinacy | The ability to cope with unpredictable events without crashing. |

1.4 Functions of an operating system

An operating system (OS) has two main complementary functions:

- resource managing; and
- machine extending.

Resource managing

It manages resources shared among users and user programs and maximises their utilisation of the CPU, RAM and other resources. This is done simultaneously in order to increase the availability.

This is similar to the role of computer operators in the 1950s.

Machine extending

It presents a virtual machine (or extended machine) to users that is much easier to access than the underlying physical machine.



The virtual machine presented to the user provides an abstraction of the computer system. This hides the complexity of the hardware from the user; this means that the user need only be concerned with the details of the hardware if they desire.

This is a way of translating the functions needed by a user from the hardware to a presentable and user-friendly medium. As a result, the operating system (OS) acts as an intermediary layer between the user and machine language.

The benefit of this abstraction can be demonstrated when comparing how computations may be processed with and without an operating system (OS).

| Without Operating System (OS) | With Operating System (OS) |
|---|--|
| <p>The instructions written in machine code or assembly language much interface directly with memory hardware. As such, the memory locations to load the two numbers from must be explicitly defined and the memory location to which the result is stored must also be defined.</p> <p>The example below shows a possible assembly code implementation of a computation that is capable of adding two numbers.</p> | <p>The instructions can be written in a high-level language, such as C++.</p> |
| LDAA \$80 LDAB \$81 ADDB STAA &55 | (load number at memory location 80) (load number at memory location 81) (add these two numbers) (store the sum to memory location 55) |
| | <pre>int a, b, c; a = 1; b = 2; c = a + b;</pre> |

Figure 1.9: Adding two numbers

This demonstrates that program development is much more user-friendly with an oper-

ating system (OS). This is because without an operating system (OS) the user must have knowledge of the system hardware; in this case, the necessary memory locations.

In addition, it is possible that the machine code or assembly language written may not work on another computer system as that computer system may have a different architecture or the memory locations may be different. For example, in another computer system:

- memory location 81 may not exist as the memory is smaller; or
- memory location 55 contains important data or instructions that should not be overwritten and therefore the computer system may crash.

By contrast, with an operating system (OS), it is possible to perform computations without interfacing directly with hardware. In the example above, variables (a, b and c) can be used to access data in memory rather than addressing memory locations. The only concern here is that the variable c is able to store the result of a + b.

The operating system (OS) provides a unified environment to users to run their computations in different systems. The operating system (OS) is capable of taking high-level code and translating it in to the machine code that can be executed on a particular computer system.

1.5 Current operating system (OS) trends

Hardware evolution

Due to fast rate of hardware evolution, operating systems (OSs) are more wide-spread than just traditional desktop computers. They can be found on hardware such as:

- mobile devices, such as smartphones and tablets; and
- embedded systems.

| Desktop | Mobile | Embedded |
|---|--|--|
| <ul style="list-style-type: none">• Windows• macOS• Linux | <ul style="list-style-type: none">• iOS• Android• Symbian OS | <ul style="list-style-type: none">• Windows Embedded |

Figure 1.10: Popular operating systems (OSs)

Multiprocessor systems

Definiton

A **multiple processor computer system** makes use of two or more processors and has the ability to allocate tasks between them.

Workstations

A single machine may contain multiple processors and therefore have large computing power.

Distributed and network systems

These computer systems share computing power and peripherals.



The diagram shows an abstraction of a distributed or network system. P1 and P2 are the connected computer systems that both share memory, access to the disk control and, if the system is a network system, the network interface.

However, there is a distinguishable difference between distributed and network systems.

| Similarities | Differences |
|--|--|
| Consist of multiple systems that are interconnected to exchange information. | In distributed systems, users are not aware of the multiplicity of computer systems available. |
| | In network systems, users explicitly move/share files, submit jobs for processing and other perform other similar tasks. |
| | In distributed systems, tasks such as moving/sharing files, submitting jobs for processing and other similar tasks are handled automatically by the operating system (OS). |

Evaluation

| Advantages | Disadvantages |
|--|--|
| Increase processor throughput due to the use of parallel processing. | A more complex operating system is required in order to be able to interface and manage multiple processor units. |
| Lower cost than using multiple processors across multiple computer systems because the processors share resources such as the power supply and motherboard. | |
| Increased reliability because failure of one processor does not affect the other processors and will only slow down the computer system. | |

1.6 Operating system (OS) layers

User interfaces

In an operating system (OS), the top layer is the user interface. This is the only layer explicitly visible to the user.

The user interface may be a:

- terminal
 - graphical user interface (GUI)
- text-based command prompt; and/or
 - a visual way of interacting with a computer using items such as windows, icons and menus.

History of the graphical user interface (GUI)

The first company to develop a graphical user interface (GUI) was Xerox PARC. They developed the “Alto” personal computer. It had a bitmapped screen and was the first computer to have a “desktop” screen with a graphical user interface (GUI).

The “Alto” personal computer was not a commercial product. However, several thousand units were manufactured and used at Xerox’s offices and several universities.

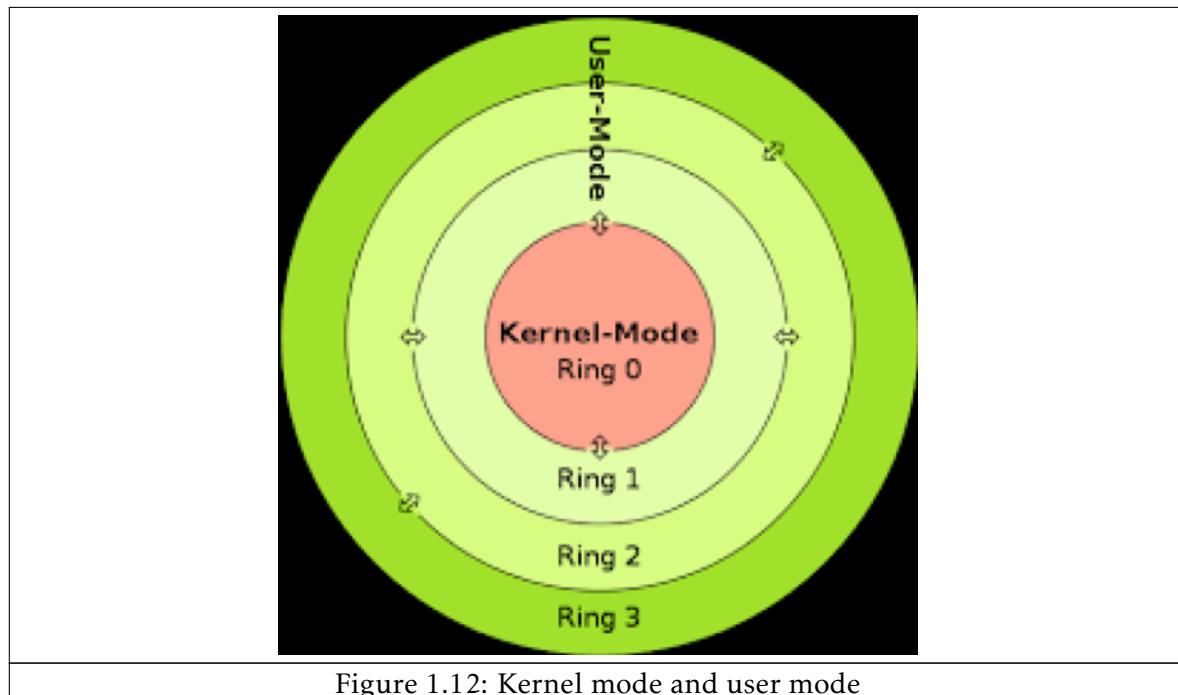
This development was a large influence on the design of personal computers during the late 1970s and early 1980s. Notable examples include:

- Three Rivers PERQ;
- Apple Lisa;
- Apple Macintosh; and
- the first Sun Workstations.

1.7 Kernel mode

Kernel mode vs user mode

| Kernel Mode | User Mode |
|--|--|
| <p>Operating systems (OSs) run in kernel mode.</p> <p>This allows:</p> <ul style="list-style-type: none"> • execution of privileged machine instructions; and • complete access and control of all the hardware. | <p>Other software runs in user mode.</p> <p>In this mode, instructions that affect control of the machine are forbidden.</p> <p>For example:</p> <ul style="list-style-type: none"> • web browsers; • e-mail software; and • music players. |



Ring 0 represents the kernel mode.

Rings 1-3 represent the user mode.

Kernel mode protection

These rings allow separation between the operating system (OS) and user programs for security and protection purposes.

If user mode had unrestricted access to all of the machine instructions:

- a user could inadvertently obtain a virus or write code that is capable of causing damage to the system, and therefore it is necessary to prevent any instructions from directly controlling the machine;
- a program may use resources unfairly, such as holding the CPU or memory, and therefore harm the execution of other programs; and/or
- sensitive machine instructions could be used improperly which may lead to kernel mode errors.

Kernel mode errors are catastrophic.

| Kernel Mode | User Mode |
|--|--|
| <p>A kernel panic represents the operating system (OS) attempting to prevent software causing any harm to the computer system and to recover from a kernel mode error on reboot.</p> <p>An example of a kernel panic is the “Blue Screen of Death” (BSOD) in Microsoft’s Windows.</p>  | <p>Application errors where an exception was thrown due to an attempt to execute a privileged instruction, this is one that should only be accessed and executed by the kernel mode, represents the operating system (OS) preventing an application from having unrestricted access to all of the machine instructions.</p>  |

Figure 1.13: Kernel mode errors

System calls

A **system call** is the programmatic way in which a computer program requests a service from the kernel of the operating system on which it is executed. A system call is a way for programs to interact with the operating system.

Some privileged instructions can be called by a programmer via system calls. Operating systems (OSs) contain system calls for which provide a layer of services for user programs to implement some activities/request services. These are usually sensitive or privileged from the kernel.

All interactions with the hardware are implemented via system calls. For example, a system call may occur if an application requires interaction with a peripheral such as a printer.

Invoking a system call is similar to calling a general function. However, the difference is that a general function's code is part of program itself, while the system call code is part of the operating system (OS). Different operating systems (OSs) offer different (limited) sets of system calls.

| Call | Description |
|-------------------------------|---|
| Process Management | |
| pid = fork() | Create a child process identical to the parent. |
| exit(status) | Terminate process execution and return status. |
| File Management | |
| n = read(fd, buffer, nbytes) | Read data from a file in to a buffer. |
| n = write(fd, buffer, nbytes) | Write data to a file |
| Process Management | |
| seconds = time(&seconds) | Get the elapsed time since Jan 1. 1970 |

Figure 1.14: Unix system calls

2. Process Management

| | | |
|-----|------------------------|----|
| 2.1 | Programs and processes | 19 |
| 2.2 | Process life cycle | 20 |
| 2.3 | Program execution | 26 |
| 2.4 | Concurrency | 28 |
| 2.5 | Process scheduling | 29 |

2.1 Programs and processes

Definitions

A **program** is the code written by a programmer.

A **process** (or job/task) shows a program in execution and is a particular instance of a program. These may be shown in a monitoring software, such as Windows Task Manager.

Data are stored values used for the computations by the process.

How it works

A single program may have multiple processes that are currently running.

Each process can share the same code for the program. This is possible as each process uses its own address space, a list of memory locations which the process can read and write. These memory locations contain the program's code instructions and data.

A program is only code however, once it is run, a process is started, and it becomes instructions and data.



2.2 Process life cycle

Process creation

A process can be created by:

- a user – a program may be executed by a user, such as via a double-click using a graphical user interface (GUI) or by typing in a command, and “trigger” the processor to load the program’s executable file containing the program code; or
- another process – an existing process may create another process by spawning/forking – the process that creates a new process is called the parent process while the new process is called the child process. The child process may also spawn a new process forming a tree of processes, as demonstrated in the diagram below.



Process table and process control block

A **process identification number (PID)** is a unique identifier given to a new process when it is created.

A **process control block (PCB)** holds all of the information about a process. It is created when new process is created.

A **process table** stores the process identification numbers (PIDs) for a process and a pointer to the respective process control block (PCB) for that process. This is managed by the operating system (OS).



Figure 2.3: Process table with respective process control blocks

The process descriptor fields in the process control block (PCB) may differ between operating system (OS). An example of possible process descriptor fields is shown by those used by the Minix operating system (OS) are shown below.

| Process Management | Memory Management | File Management |
|---------------------------|-------------------------|------------------------|
| Registers | Pointer to text segment | UMASK mask |
| Program counter | Pointer to data segment | Root directory |
| Program status word | Pointer to bss segment | Working directory |
| Stack pointer | Exit status | File descriptors |
| Process state | Signal status | Effective uid |
| Time when process started | Process ID | Effective gid |
| CPU time used | Parent process | System call parameters |
| Children's CPU time | Process group | Various flag bits |
| Time of next alarm | Real uid | |
| Message queue pointers | Effective uid | |
| Pending signal bits | Real gid | |
| Process ID | Effective gid | |
| Various flag bits | Various flag bits | |

Figure 2.4: Process descriptor fields in Minix

Although Minix is a fully-featured operating system (OS), it is a small operating system

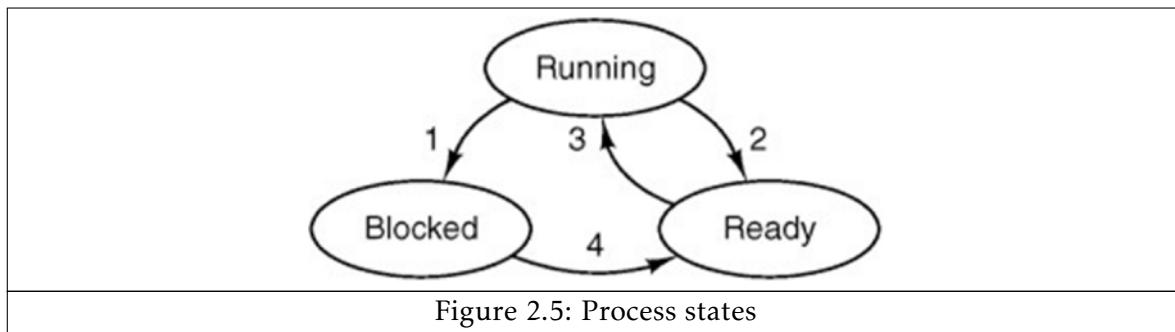
(OS) and therefore the processor descriptor fields are less complex than other operating systems (OSs).

Three-state model

The **three-state model** shows how a process, once initiated, can be in one of three states. The current state of a process is stored in its respective process control block (PCB).

A process, once initiated, can be in one of three main states:

- running – actually using the CPU to perform a task;
- ready – ready to run but waiting for the CPU as it has not yet had time on the CPU has been temporarily stopped to let another process run; or
- ready – unable to run until some external event occurs, such as:
 - waiting for an interrupt, this is a message from the hardware saying that a resource is now available to read from – such as waiting for an input/output (I/O) operation to complete; and
 - waiting for another process to finish accessing a shared resources – for example: a file; memory; or an external peripheral, such as a printer.



In reference to the diagram above, transitions can occur between process states.

| Transition | Diagram Number | Explanation |
|--------------------------|----------------|---|
| running → blocked | 1 | Process blocks for input. For example, if the process is waiting for some input from I/O. |
| running → ready | 2 | Scheduler picks another process. The process has had opportunity to run and is flagged as no longer currently running, so that another process can run. |
| ready → running | 3 | Scheduler picks this process. The next process that is ready to run is set to running to allow access to the CPU. |
| blocked → ready | 4 | Input becomes available. The process has received input from I/O or the process sends an interrupt and the interrupt service routine (ISR) is executed, the scheduler is called to transition the process from blocked to ready. |

Figure 2.6: Transitions between processor states

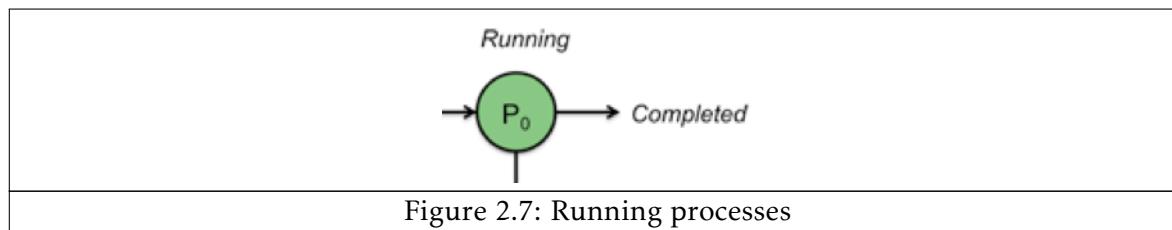
The transition between process states is dependent on the scheduling algorithm used. A clock is used to send a signal to stop the current process, move it from running to ready and then run the scheduler to find out what process should be processed next and the next process will be made to running.

State queues

At any time, a process is in only one state.

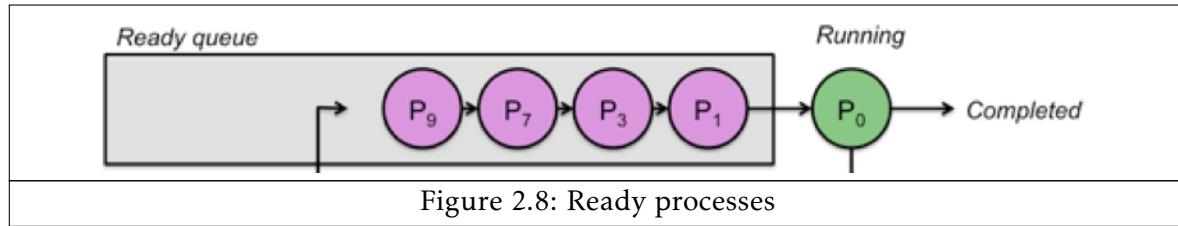
Running processes

At any time, at most one process is in the running state. This is because a single-core processor is only able to process one instruction coming from a single core at a time. If a computer system has a multi-core processor, this rule applies to each individual core on the processor, rather than the processor as a whole, as they are able to complete parallel execution.



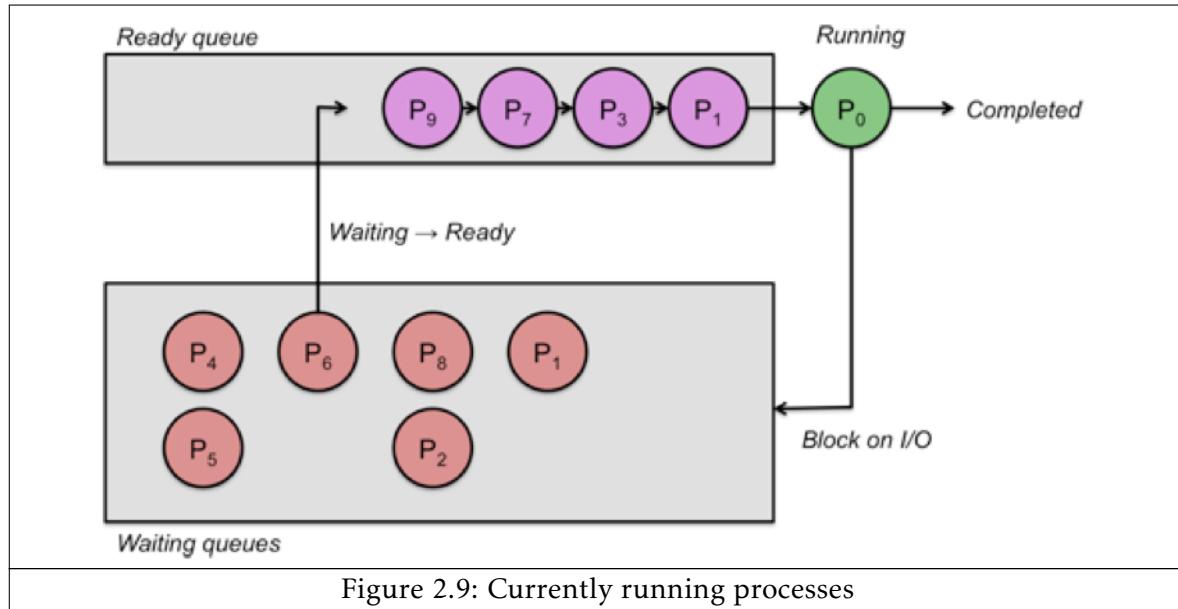
Ready processes

There may be a queue of processes in the ready state.



Blocked processes

There may be several queues of processes in the blocked state, where each queue represents one resource for which processes in that queue are waiting.



Process termination

Process termination is the end of life for a process, this can occur in two general ways.

Voluntary termination

Voluntary termination of a process represents the end of life for a process where its termination was intended by the user or the programmer.

This can be a:

- **normal exit** – the process has done its work; or

- error exit**
 - the process itself handles and “catches” an error – for example, some try-catch code is implemented to check if a condition is met, such as if the definition of a variable is present.



Figure 2.10: Voluntary termination due to an error exit

Involuntary termination

Involuntary termination of a process represents the end of life for a process where its termination is not intended by the user or the programmer.

This can be due to:

- a fatal error**
 - an error is detected by the operating system's (OS's) protected mode – for example, an exception has been thrown due to reference to a non-existent memory location or division by zero; or
- being killed by another process**
 - a process may execute a system call that causes the operating system (OS) to kill another process – this process may have control over the killed process and this may be due to that process being the parent process of the killed process.

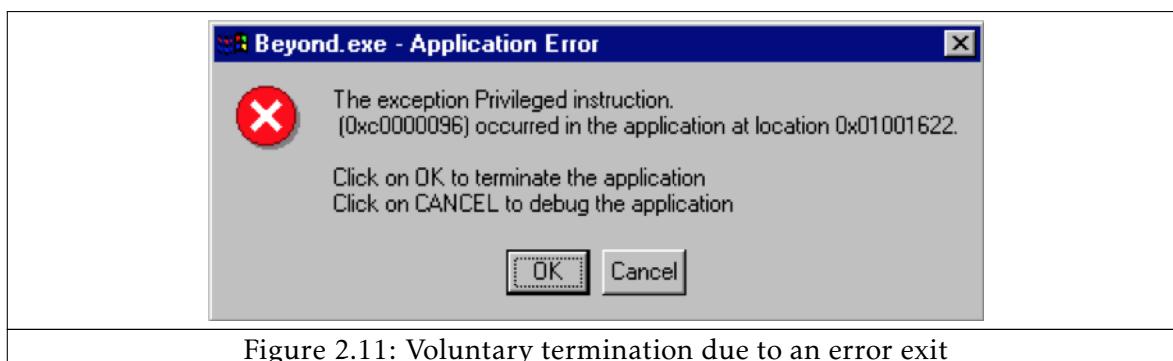


Figure 2.11: Voluntary termination due to an error exit

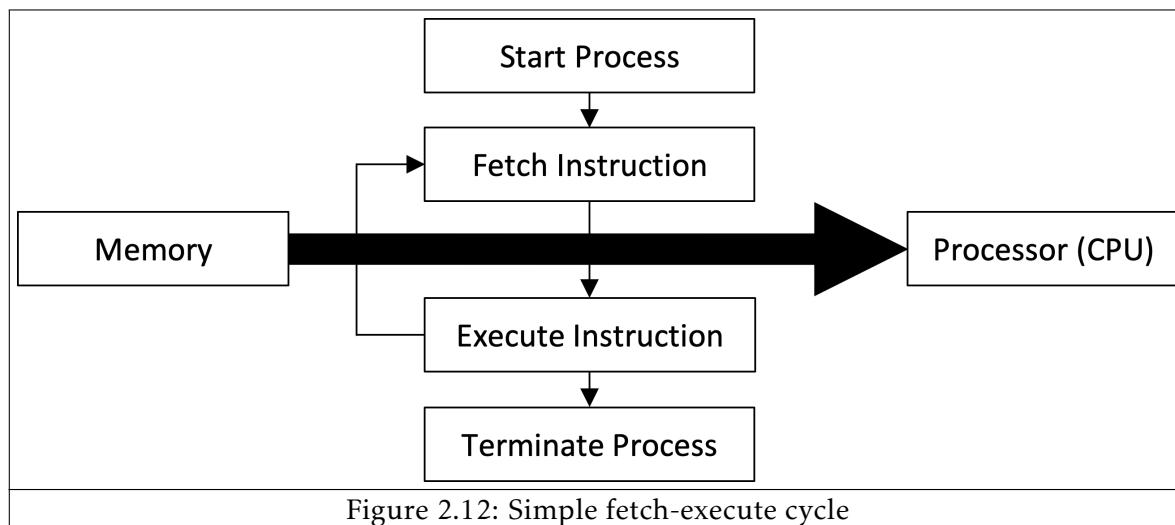
2.3 Program execution

The simple fetch-execute cycle

Definition

The **fetch-execute cycle** is an operational process in which a computer system retrieves a program instruction from its memory, determines what actions the instruction dictates, and carries out those actions.

How it works



Once a process is started, instructions are fetched from memory and executed using the CPU. This process will continue until the process is terminated.

This predictable cycle is only feasible to an extent as some processes may be slow or blocked and others may require immediate attention and cannot wait for the current process to terminate. For example, if a process blocks the processor (CPU) because it is waiting for an event to occur, such as a printer to finish its job, or if a high-priority process is supposed to execute as soon as possible. In these cases, interrupts are required.

Interrupts

Definition

An **interrupt** is a signal sent to the processor indicating that an event caused by hardware or software requires immediate attention.

Types of interrupts

- **Input/output (I/O) interrupt** – Caused by an input/output (I/O) device to signal completion or an error.
- **Timer interrupt** – Caused by a processor timer and is used to alert the operating system (OS) at specific instants.
- **Program interrupt** – Caused by error conditions within user programs or fatal errors.

When do interrupts occur?



Interrupts enable operating systems (OSs) to oversee several programs and input/output (I/O) events simultaneously.

This also means that single-core processors can effectively emulate the way in which multi-core processors deal with multiple instructions at a given time by switching between instructions intelligently. Due to the high clock speeds of modern processors, it is easy to give the illusion that true multi-tasking, where two instructions are being processed at once, is taking place on a single-core processor.

Updated fetch-execute cycle

Given the introduction of interrupts, it is now necessary to update the fetch-execute cycle in order to include the possibility of an interrupt.

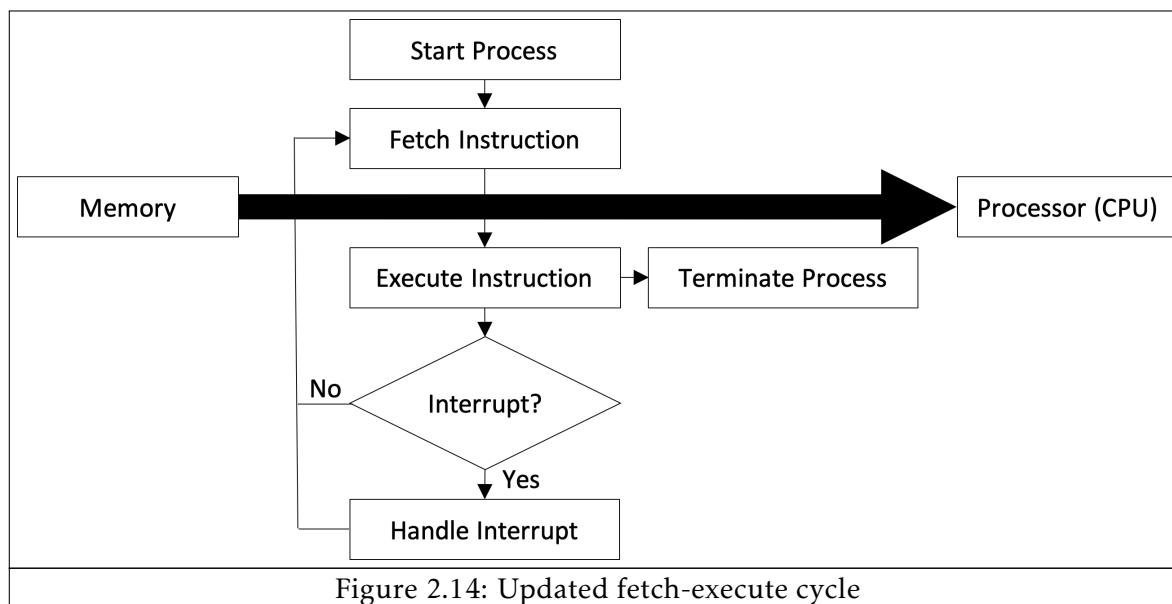


Figure 2.14: Updated fetch-execute cycle

2.4 Concurrency

Definition

Concurrency describes the ability for a program to be decomposed into parts that can run independently of each other. This means that tasks can be executed out of order and the result would still be the same as if they are executed in order.

Why is concurrency required

Concurrency allows the processor (CPU) to run several processes. An example of this is shown by an interrupt occurring due to an input/output (I/O) event occurring (page 24).

Interleaving

Concurrency is able to achieve multitasking, that does not require parallel execution, by performing interleaved execution.

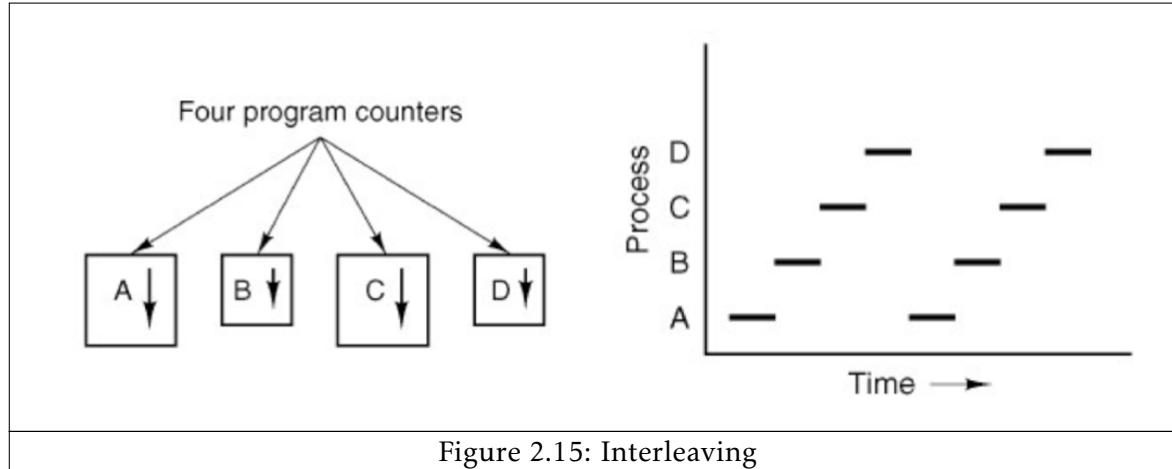


Figure 2.15: Interleaving

2.5 Process scheduling

The scheduler

Definition

A **scheduler** uses a scheduling algorithm to determine how to share processor time.

How it works

After an input/output (I/O) system call or interrupt handling, control is passed to the scheduler to decide which process to execute next.



The scheduler checks if the current process is still the most suitable to run at this moment in time. If it is control is returned to the process, otherwise:

- the state of the current process is saved in the process control block (PCB);
- the state of the most suitable process is retrieved from the process control block (PCB); and
- control is transferred to the newly selected process at the point indicated by the restored program counter (PC).

The action of storing the state of the current process and activating another process is called a context switch.

Context switching must be minimised to reduce overheads created by copying the state of processes and the time taken to perform the switch. However, it is still regarded as important to context switch when appropriate to avoid longer waiting times in the event of a blocked process.

Scheduling

Definition

Scheduling is the act of determining the optimal sequence and timing of assigning execution to processes.

Scheduling criteria

Different scheduling criteria may be selected depending on the use case for a given computer system.

- **CPU utilisation** – Aims to keep the CPU as busy as possible.
- **Efficiency** – Aims to maximise system throughput.
- **Fairness** – Aims to be fair to all running processes or to all users on a multi-user operating system (OS).

This means that different policies and algorithms for scheduling will exist to match these criteria.

Scheduling policies

A **non-preemptive scheduling policy** is one that allows processes to run until complete or incurring an input/output (I/O) wait. These scheduling policies can be described as passive.

A **preemptive scheduling policy** is one that allows processes to be interrupted and replaced by other processes, generally through timer interrupts.

Scheduling algorithms

Definitions

Arrival time is the instant at which a process is first created.

Service time is the time that it takes for a process to complete if it is in continuous execution.

The **waiting time** for a process is the sum of time spent in the ready queue during the life of the process. This does not include time that the process is blocked or waiting for input/output (I/O).

In order for a scheduling algorithm to be deemed as fair to all processes, the ratio between waiting time and run time should be about the same for each process.

Case study

| Process | Arrival Time | Service Time |
|---------|--------------|--------------|
| A | 0 | 3 |
| B | 2 | 6 |
| C | 4 | 4 |
| D | 6 | 5 |
| E | 8 | 2 |

Figure 2.17: Case study

The case study shows different processes (A-E) that all have different arrival times and different service times.

The following examples of scheduling algorithms will refer to this case study to demonstrate their function.

First come, first served (FCFS) / First in, first out (FIFO) – Non-preemptive

In this algorithm, the first process to arrive is assigned to the processor (CPU) until it is finished. Meanwhile, any other processes that come along are queued up waiting to be processed.



| Advantages | Disadvantages |
|----------------------|---|
| Simple to implement. | Does not consider the priority of a process and therefore the important processes may not be completed quickly. |

| | |
|---|---|
| | Prevents other processes from starting while another is in progress and therefore, if processes are of varying sizes, there may be inefficiencies since a single process may take a long time to complete therefore leaving the user waiting before they can perform any other actions. |
| Favours long processes as all processes are given the opportunity to run until completion and therefore, some shorter processes may not be able to start processing until the longer processes are finished. | |

Shortest job first (SJF) – Non-preemptive

In this algorithm, the process with the shortest estimated run time is assigned to the processor (CPU) until it is finished. Meanwhile, any other processes that come along are queued up waiting to be processed.

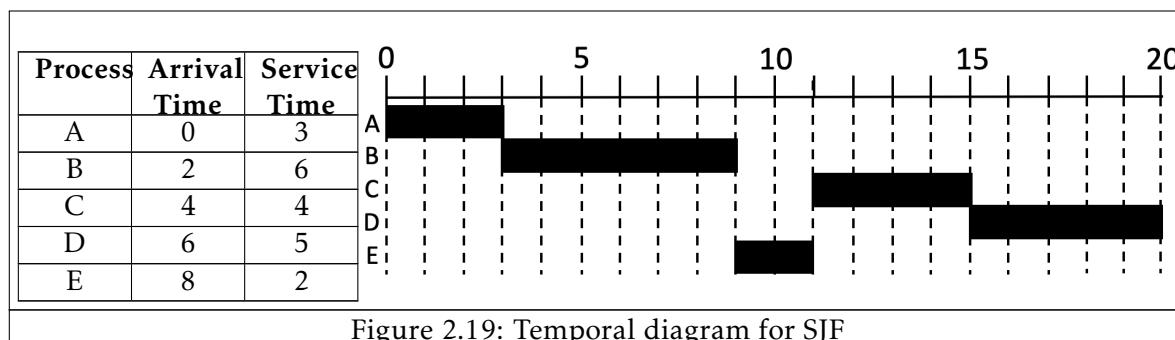


Figure 2.19: Temporal diagram for SJF

| Advantages | Disadvantages |
|--|--|
| Simple to implement. | Does not consider the priority of a process and therefore the important processes may not be completed quickly. |
| Shorter processes are processed quickly because they take precedence. | Relies on an estimation of how long a process will take which could be incorrect. |
| Minimises the average time taken to complete a process because the shortest processes take precedence. | Prevents other processes from starting while another is in progress and therefore, if processes are of varying sizes, there may be inefficiencies since a single process may take a long time to complete therefore leaving the user waiting before they can perform any other actions. |
| Favours long processes as the processes with the shortest estimated time are given the opportunity to run until completion and therefore, some longer processes may not be able to start processing until there are no more shorter processes currently running or ready to be processed. | |

Shortest remaining time (SRT) – Preemptive

In this algorithm, the process with the shortest estimated remaining time is assigned to the processor (CPU). If a process becomes ready that has a shorter remaining time, it will be pre-empted to allow the new process to start and the scheduler is switched to that new process.



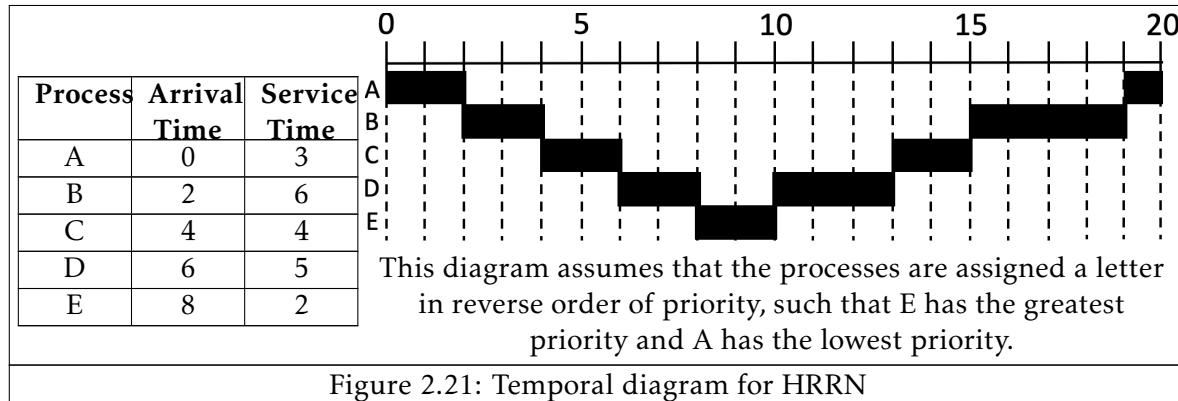
| Advantages | Disadvantages |
|--|--|
| High throughput because the number of processes completed is high due to the shortest processes taking precedence. | Does not consider the priority of a process and therefore the important processes may not be completed quickly. |
| Shorter processes are processed quickly because they take precedence. | Relies on an estimation of how long a process will take which could be incorrect. |
| Favours long processes as the processes with the shortest estimated time are given the opportunity to run until a process with a shorter estimated time is ready and therefore, some longer processes may not be able to start processing until there are no more shorter processes currently running or ready to be processed. | |

Highest response ratio next (HRRN) – Preemptive

In this algorithm, the process with the highest priority is assigned to the processor (CPU). If a process becomes ready that has a higher priority, it will be pre-empted to allow the new process to start and the scheduler is switched to that new process.

The priority of a process may be based on memory, time and/or any other resource requirement such as:

$$\frac{\text{waiting time} + \text{runtime}}{\text{runtime}}$$



| Advantages | Disadvantages |
|---|---|
| High throughput because the number of processes completed is high due to the shortest processes taking precedence. | Does not consider the priority of a process and therefore the important processes may not be completed quickly. |
| Shorter processes are processed quickly because they take precedence. | Relies on an estimation of how long a process will take which could be incorrect. |
| | Can be inefficient if a large process is in progress and shorter processes are being added to the queue because they will take precedence. |

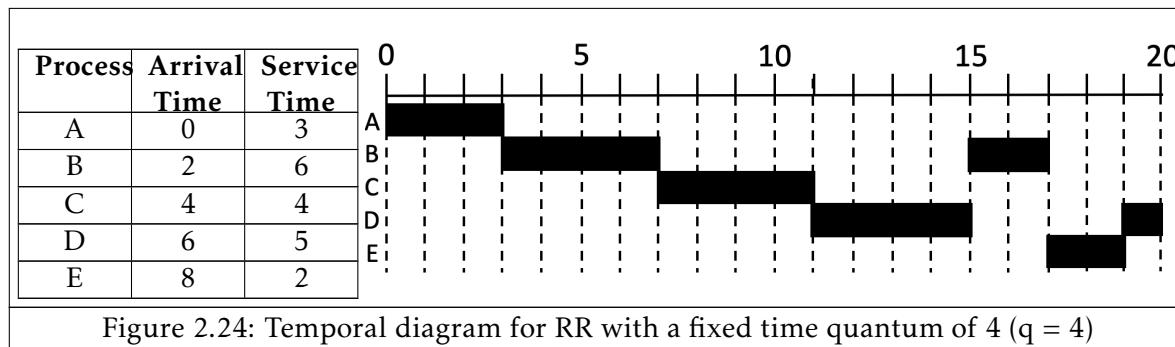
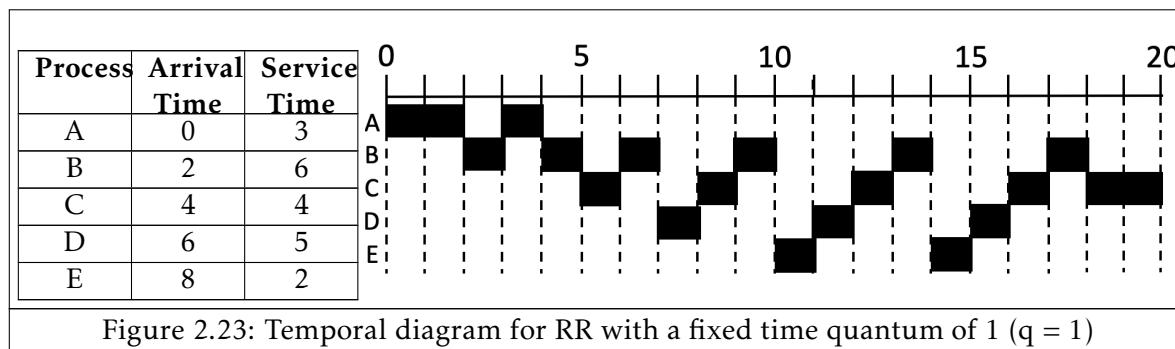
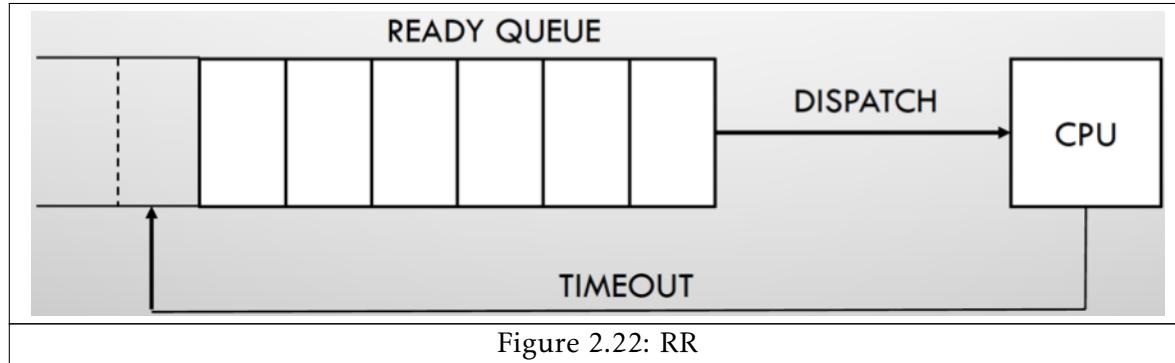
Round robin (RR) – Preemptive

In this algorithm, each process is dispatched to the processor (CPU) on a “first in, first out” (FIFO) basis with a fixed time quantum.

Each time quantum is typically 10-20ms. Modern processor (CPU) clock frequencies are typically greater than 2GHz, which implies clock periods of 5×10^{-7} ms. This shows that the given time quantum is relatively large compared to a typical processor’s (CPU’s) clock period.

If a process experiences a timeout, this will mean that it has run over its fixed time quantum. In which case, the process will be interrupted and returned to the back of the queue.

A system designer may wish to choose a time quantum that is most appropriate for a given system. This may be done by measuring the average service time and waiting time for the processes that will be running on the system and design the fixed time quantum around these figures. It may be that a system designer wishes to minimise the number of processes that are interrupted by another process.



| Advantages | Disadvantages |
|--|---|
| Simple to implement. | Heavy overhead due to continuous context switches. |
| Suitable for some types of computer systems , such as those which will be running processes of similar priority and size. | Does not consider the priority of a process and therefore the important processes may not be completed quickly. |
| | Does not consider the size of a process and therefore, if processes are of varying sizes, there may be inefficiencies since a single process may take a long time to complete thus leaving the user waiting before they can perform any other actions. |

Multi-level queueing

Definition

Multi-level queueing is a queue with a predefined number of levels which consist of several independent queues.

How it works

Multi-level queueing makes use of other existing scheduling algorithms.



Figure 2.25: Multi-level queueing

Each queue has a different priority; the top queue has the highest priority and the bottom queue has the lowest priority. Processes are able to move between queues if their priority changes.

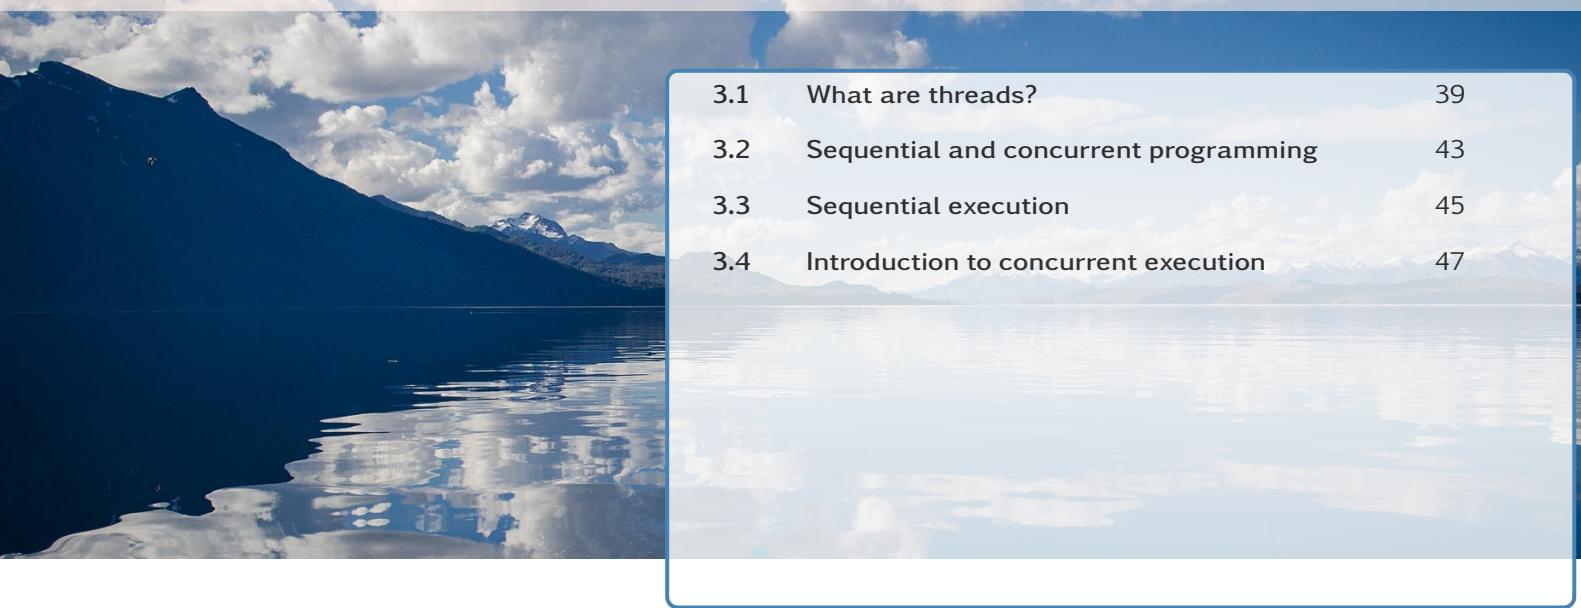
There is some form of inter-queue scheduling policy that governs the assignment of processes from each queue to the processor (CPU).

The queues in a multi-level queueing system may differ between operating system (OS). For example, the scheduler used by the Minix operating system (OS) uses multi-level queueing and implements 16 queues.

| Advantages | Disadvantages |
|---|---------------|
| Allows scheduling optimisation as a system designer may be able to leverage the advantages of a range of different scheduling algorithms. | |
| Maintains common processes as it is possible to split processes into different queues depending on their nature. For example, input/output processes could be in one queue while processor (CPU) processes are in another queue. | |

| | |
|--|--|
| Helps to prevent bottlenecks because input/output (I/O) devices are slower than the processor speed and therefore maximising processes involving input/output (I/O) devices ensures that these devices are continuously busy. | |
|--|--|

3. Threads and Concurrency



| | | |
|-----|---------------------------------------|----|
| 3.1 | What are threads? | 39 |
| 3.2 | Sequential and concurrent programming | 43 |
| 3.3 | Sequential execution | 45 |
| 3.4 | Introduction to concurrent execution | 47 |

3.1 What are threads?

Definition

A **thread** is an independent path/sequence of execution within a process and can be managed independently by a scheduler. A process may contain many threads.

Multi-threading

How it works

As mentioned in the previous section, a process (or job/task) shows a program in execution and is a particular instance of a program. By default, these processes are run by means of “single execution thread”.

However, different parts of the same process could be parallelised in order to allow multi-threading.

Multi-threading provides a way of improving application performance and therefore improving the efficiency and/or usability of a computer system.



Figure 3.1: Single-threaded vs multi-threaded

Example

A screenshot of Microsoft Word's 'Document1 - Microsoft Word' window. The main document area contains two paragraphs of text. Below the document, there is a 'Word Processor Checklist' section with links to 'Jump to a topic in this article' and 'History of Word Processing' and 'Standard Features of Word Processors'. The ribbon menu at the top includes 'File', 'Home', 'Insert', 'Page Layout', 'References', 'Mailings', 'Review', and 'View'. The 'Home' tab is selected, showing various font and paragraph styling tools. In the bottom right corner of the window, there is a status bar with the text 'In this example, a thread could be assigned to each of the following tasks:' followed by a bulleted list of tasks.

In this example, a thread could be assigned to each of the following tasks:

- user input, such as keyboard and mouse input;
- auto-saving document;
- spell-checking document; and
- printing in the background.

This is possible as all of these tasks can be run in parallel as they do not require information from one another.

Figure 3.2: Example of multi-threading in Microsoft Word

Threads vs processes

Comparison

Threads and processes share some similarities however, there are also some distinct differences.



Figure 3.3: Single-threaded process



Figure 3.4: Multi-threaded process

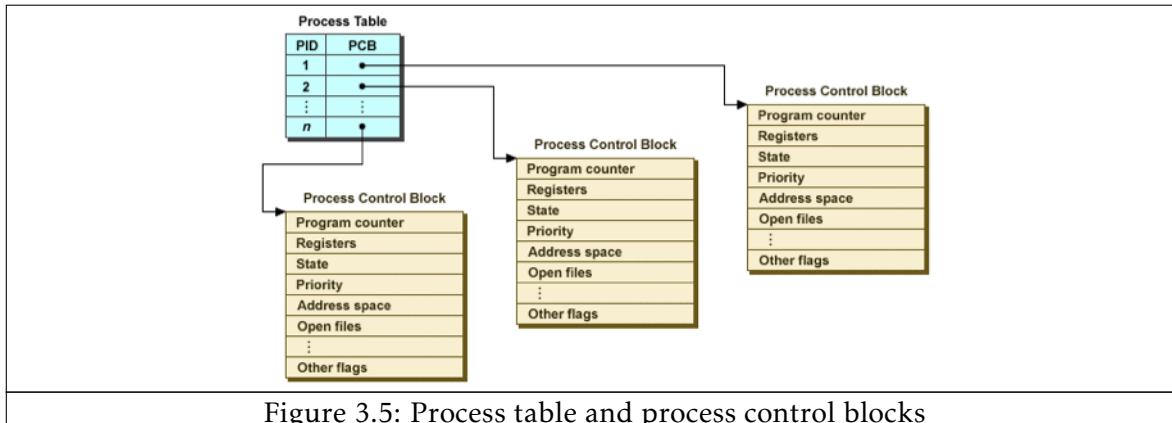


Figure 3.5: Process table and process control blocks

| | Threads | Processes |
|--------------|--|---|
| Similarities | Sequential flow of control with start and end. | Sequential flow of control with start and end. |
| | At any time, a thread has a single point of execution. | At any time, a process has a single point of execution. |
| | Has its own execution context, stack (history) and program counter stored in a thread control block (TCB). | Has its own execution context, stack (history) and program counter stored in a process control block (PCB). |
| | Follows the three-state model in which the thread can be running, blocked or ready. | Follows the three-state model in which the process can be running, blocked or ready. |
| | Context switching can happen for threads. | Context switching can happen for processes. |
| | A thread can spawn another thread. | A process can spawn another process. |
| | A thread is often called a lightweight process. | |
| Differences | A thread cannot exist on its own, instead it exists within a process. | A process does not require a parent entity. |
| | Usually created and/or controlled by a process. | A process is not typically created and/or controlled by another process. |
| | Threads can share process properties, including memory and open files. | Processes cannot share process properties with other processes. |
| | Inexpensive creation and context switching as does not require separate address space. | Expensive creation and context switching as requires separate address space. |
| | When running multiple threads concurrently, they share an address space. | When running multiple processes concurrently, they are resources, such as memory, disk and printers. |

Figure 3.6: Threads vs processes

Properties

| Per process items | Per thread items |
|-----------------------------|------------------|
| Address space | Program counter |
| Global variables | Registers |
| Open files | Stack |
| Child processes | State |
| Pending alarms | |
| Signals and signal handlers | |
| Accounting information | |

Figure 3.7: Threads vs processes

Process properties are shared between threads. Thread properties are local and private to each thread.

3.2 Sequential and concurrent programming

Sequential programming

Definition

Sequential programming is the traditional activity of constructing a computer program using a sequential programming language.

How it works

This involves a programming methodology that assumes statements are executed in order/sequence.

Programs written using sequential programming are assumed to execute on a single-CPU system and have a single thread of control.



Evaluation

| Advantages | Disadvantages |
|---|---|
| No additional support required from the programming language. | Lower processor throughput than concurrent programming as it cannot benefit from multitasking or concurrent processing. |
| No additional support required from the operating system as most old-school operating systems were generally single-threaded and therefore later generations of operating systems typically inherit this functionality. | Multiple computer systems that each have their own CPU may yield a higher cost than multi-CPU systems as more will be spent on resources, such as the power supply and motherboard, that could be otherwise shared in a multi-CPU system. |
| | Lower reliability than a multi-CPU system because, in the case of the failure of the processor, there is no redundancy. |

Concurrent programming

Definition

Concurrent programming is the activity of constructing a computer program that takes advantage of concurrency allowed by the use of multiple threads of control.

How it works

Multiple threads of control allow a given process to perform multiple computations in parallel and to control simultaneous external activities.

The program may be run on both:

- a single-CPU – where the computer program will take advantage of multitasking; system and
- a multi-CPU – where the computer program will take advantage of true system parallelism.

Evaluation

| Advantages | Disadvantages |
|---|--|
| Increase processor throughput due to the use of multitasking in a single-CPU system or parallel processing in a multi-CPU system. | Requires support from the programming language as it must implement techniques to deal with multitasking on a single-CPU system and/or parallel processing on a multi-CPU system. |
| A multi-CPU system generally yields a lower cost than using multiple CPUs across multiple computer systems because the processors share resources such as the power supply and motherboard. | Requires support from the operating system as it must support multi-threading in order to allow multitasking on a single-CPU system and/or interface and manage multiple CPU units on a multi-CPU system. |
| Increased reliability in a multi-CPU system because failure of one processor does not affect the other processors, instead the computer system may experience lower performance until fixed. | |

3.3 Sequential execution

Definition

Sequential execution is where the execution of threads in a sequential program is executed in sequence/order with no overlapping.

Order and precedence

Explanation

In sequential execution, there is only one possible sequence of execution. This is because a sequential program gives the system strict instructions on the order of executing the state-

ments in the program.

Importance

For example, a simple hypothetical program could be:

P;

Q;

R;

This tells the computer system that the statements must be executed in the order they are written, such that:

- P must precede Q; and
- Q must precede R.

High level

The importance of the order of precedence can be highlighted by demonstrating this idea in a high-level programming language.

Given the following program written in a high-level language:

```
x = 1;  
y = x + 1;  
x = y + 2;
```

it is possible to see that the final values of x and y depend on the order of execution of the statements.

System level

Given the following program written in a high-level language:

```
x = 1; P  
y = x + 1; Q  
x = y + 2; R
```

where each statement is assigned a letter respectively, each statement may be compiled into several machine instructions.

Statement **P** is treated as a single machine instruction:

- **P1**: store 1 at the memory address of x.

Statement **Q** is broken into three machine instructions:

- **Q1**: load the value of x into a CPU register;

- **Q2:** increment the value in this register by 1; and
- **Q3:** store the value in this register at the memory address of y.

Statement **R** is broken in to three machine instructions:

- **R1:** load the value add of y in to a CPU register;
- **R2:** increment the value in this register by 2; and
- **R3:** store the result at the memory address of x.

The nature of sequential execution

The execution of statements **P**, **Q** and **R** at the program level (or high-level) as

P → **Q** → **R**

implies that the execution at the system level is as follows

P1 → **Q1** → **Q2** → **Q3** → **R1** → **R2** → **R3**,

given that **P** is compiled to a single machine instruction, whilst **Q** and **R** are compiled to three machine instructions – as seen on page 40.

Sequential execution has the following assumptions:

- total ordering
 - there is single-threaded computation, and therefore no overlap in the execution of the statements;
- deterministic
 - the same input will always result in the same output; and
- sequence
 - users will specify a strict sequence of steps required in order to achieve a desired goal.

However, this does not apply in many practical applications, for which a sequence of steps is not required.

3.4 Introduction to concurrent execution

Definition

Concurrent execution is where the execution of threads in a concurrent program is occurring asynchronously, meaning that the order in which tasks are executed is not predetermined.

The squares example

In this hypothetical example, a person desires to have a list with the results of all of the squares (2^n) from 1 to 100000.

A group of 100000 people are split in to heavily uneven teams and assigned the same task to complete all of the calculations in order to achieve the desired result.

It is given that each calculation takes n amount of time.

| Team 1 | | Team 2 | |
|-------------------|---|-------------------|---|
| Number of members | 1 | Number of members | 99999 |
| Strategy | One person should complete all of the calculations. | Strategy | Each member is assigned a number between 1 to 100000. Each member should calculate the respective square for the number they are assigned. |
| Time taken | $100000n$ | Time taken | n |

This shows that Team 2 was $100000x$ faster than Team 1. This was because it was possible to decompose the larger task in to smaller sub-tasks and assign each of those tasks to a separate resource, which in this case is one person.

This example forms the basic concept of concurrent execution.

The nature of concurrent execution

Concurrent execution dismisses many of the assumptions required for sequential execution (page 41).

Calculations may be carried out without total ordering. As a result, calculations may be carried out in parallel and overlapping is therefore allowed.

In the example above, each individual person in team 2 carried out their operations in sequence.

In the example above, the operations in the whole computation can be viewed as being in partial order. However, the order does not matter here because there is no dependency between the calculations. This is because the output from any given calculation is not required as an input to any other given calculation.

However, in general, concurrent execution is non-deterministic, and therefore the same input generally means different output due to ordering. This is because there are many cases where the order of operation does matter.

