

COMP20081 SYSTEMS SOFTWARE

REVISION GUIDE

BY MATT ROBINSON

Copyright © 2017 Revision Guide

Copying prohibited

All rights reserved. No part of this publication may be reproduced or transmitted in any form or by any means, electronic or mechanical, including photocopying and recording, or by any information storage or retrieval system, without the prior written permission of the publisher.

Art. No xxxxx

ISBN xxx-xx-xxxx-xx-x

Edition 0.0

Cover design by Cover Designer

Published by by Matt Robinson
Printed in City

Contents



1	Introduction	5
1.1	History of performing computations	5
1.2	Hardware	7
1.3	What is an operating system (OS)?	9
1.4	Functions of an operating system	11
1.5	Current operating system (OS) trends	13
1.6	Operating system (OS) layers	15
1.7	Kernel mode	16
2	Process Management	19
2.1	Programs and processes	19
2.2	Process life cycle	20
2.3	Program execution	26
2.4	Concurrency	28
2.5	Process scheduling	29
3	Threads and Concurrency	39
3.1	What are threads?	39
3.2	Sequential and concurrent programming	43
3.3	Sequential execution	45
3.4	Introduction to concurrent execution	47

3.5	Interleaving	48
3.6	User and kernel threads	52
3.7	Multi-threading models	54
3.8	Evaluation of concurrent programming	57
4	Synchronisation and Mutual Exclusion	59
4.1	Race condition	59
4.2	Inter-process synchronisation	61
4.3	The producer-consumer problem	65
4.4	The dining philosophers problem	71
5	Memory Management	77
5.1	Memory hierarchy	77

1. Introduction

1.1	History of performing computations	5
1.2	Hardware	7
1.3	What is an operating system (OS)?	9
1.4	Functions of an operating system	11
1.5	Current operating system (OS) trends	13
1.6	Operating system (OS) layers	15
1.7	Kernel mode	16

1.1 History of performing computations

1950s

A user may want to perform computations that may have a complexity that is not feasible by a standard calculator.

A description of the computations to be performed could be expressed in a physical form using a perforated card.

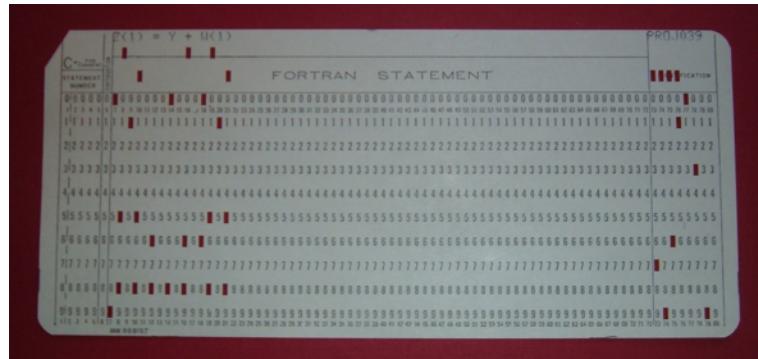


Figure 1.1: Perforated card

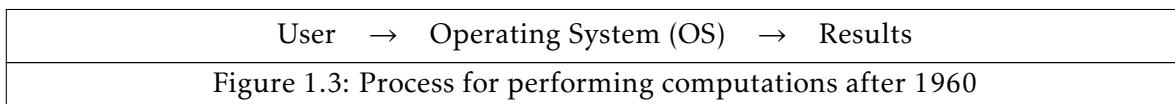
There are multiple lines on the card, each with varying holes. Each line describes a particular instruction that is necessary for the computations to be completed.



>1960s

The process of performing computations after 1960 changed dramatically and required far less human input.

Automation was brought about by the “Operating System Paradigm” in which users could interface with a computer system directly through an operating system (OS).



History details

Late 1950s	<ul style="list-style-type: none"> • Standard subroutines were produced that were loaded at start-up. These contained features similar to those found on an operating system. • Magnetic tapes were used for storage and were later replaced by disks. • Assemblers started to be used. These are programs that takes basic computer instructions and convert them in to machine code; this is a pattern of binary bits (0's and 1's) that the computer system's processor can use to perform its basic operations. • High-level languages, which consisted of more natural and human-readable language, started to be used. For example, FORTRAN is a general-purpose, compiled imperative programming language that is especially suited to numeric computation and scientific computing and was introduced in 1957.
------------	--

1960s	Automated batch system. <ul style="list-style-type: none">• This replaced the computer operator.• Several programs could be loaded into memory and automatically processed in a “first in, first out” (FIFO) fashion.
1970s	Multiprogramming. <ul style="list-style-type: none">• The computer could switch between jobs, which allows processing and input/output (I/O) interaction simultaneously.
1980s	Graphical user interfaces (GUIs). <ul style="list-style-type: none">• The interaction between a computer system and a user through the medium of a mouse and keyboard.

1.2 Hardware

External hardware

A **peripheral** is any external hardware device that provides input/output (I/O) for the computer.

For example, a keyboard and mouse are input peripherals, while a monitor and printer are output peripherals. Some peripherals, such as external hard drives, provide both input and output for the computer.

A computer system generally has many internal hardware components and hardware peripherals.



Figure 1.4: External view of computer system

Internal hardware

A **processor** or **central processing unit (CPU)** is the hardware within a computer that carries out the instructions of a computer program by performing the basic arithmetical, logical, and input/output operations of the system.

A **motherboard** is the main printed circuit board (PCB) in a computer. The motherboard is a computer's central communications backbone connectivity point, through which all components and external peripherals connect.

Inside of a computer system, there are many components connected to the processor via the motherboard.

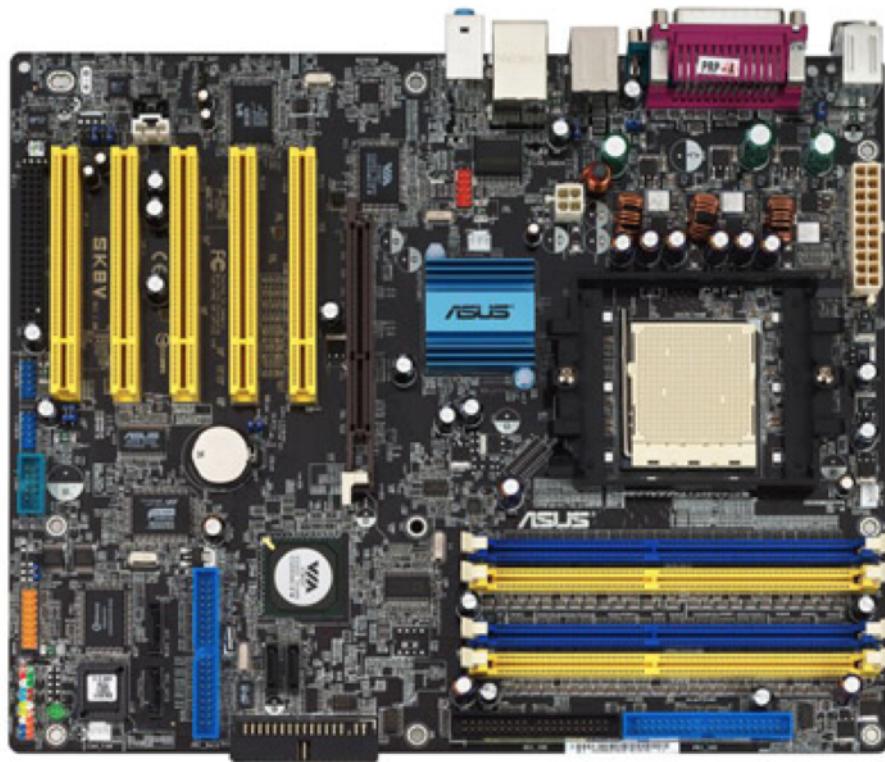


Figure 1.5: Typical computer motherboard

The motherboard connects:

- all of the internal components via the data bus; and
- the peripherals.

1.3 What is an operating system (OS)?

Definition

An **operating system (OS)** is a collection of system programs that manage the hardware resources and peripherals connected to a computer system. It is also responsible for the graphical user interface or command line interface and all other software running on the computer system.

Purpose of an operating system (OS)

An operating system (OS) is designed to:

- eliminate the need to have hardware knowledge to operate a computer system;

- make the boundary between hardware and software transparent, allowing the user to not be concerned with the technical details; and
- provide a user-friendly environment to execute and develop programs.

These attributes are achieved by layering the computer system such that the user can interface with applications, rather than the operating system (OS) or the hardware directly.

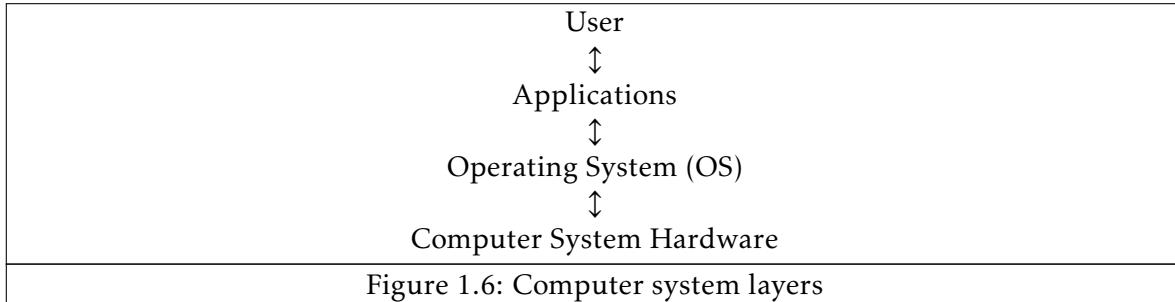


Figure 1.6: Computer system layers

Structure of an operating system (OS)

The structure of an operating system (OS) can be said to resemble an onion.

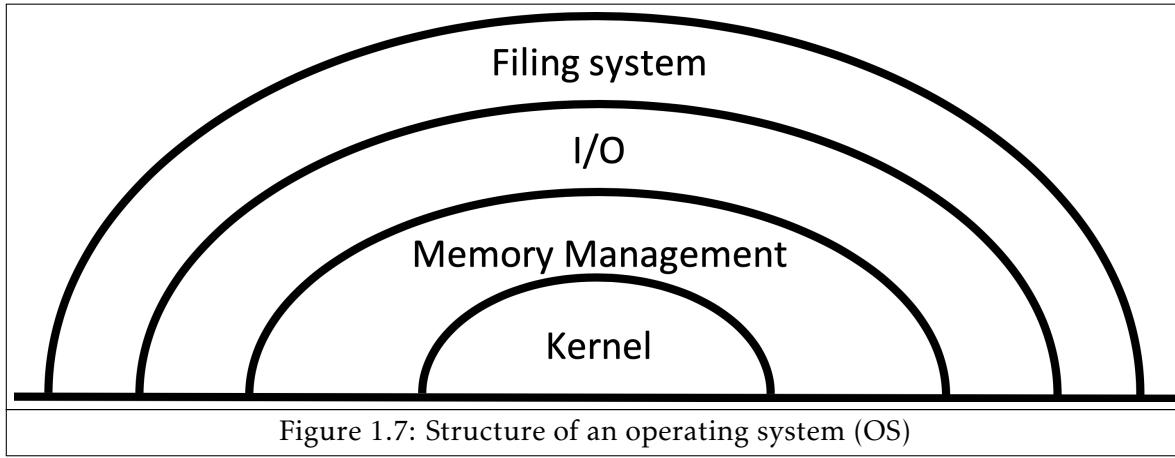


Figure 1.7: Structure of an operating system (OS)

An operating system (OS) has four main components.

The **kernel** hides the complexity of how a computer system works from users. It is responsible for:

- process management;
- CPU scheduling; and
- handling interrupts.

Memory management is responsible for allocating and deallocating memory to processes.

Input/output (I/O) includes any interaction between the internal computer system components and peripherals.

The **filing system** is comprised of file management subsystems.

Each layer in the operating system (OS) structure provides functions to the above layers.
Each layer uses facilities provided by layers within and below that layer.

Practical features of current operating systems (OSs)

Concurrency	Allows overlapping input/output (I/O) operations with computations and several programs to be stored in memory at a single time.
Sharing of resources	Sharing hardware and peripherals, such as hard disks and printers.
Access to long term storage	Important for saving important files on mediums such as hard disk drives (HDDs) and solid-state drives (SSDs).
Non-determinacy	The ability to cope with unpredictable events without crashing.

1.4 Functions of an operating system

An operating system (OS) has two main complementary functions:

- resource managing; and
- machine extending.

Resource managing

It manages resources shared among users and user programs and maximises their utilisation of the CPU, RAM and other resources. This is done simultaneously in order to increase the availability.

This is similar to the role of computer operators in the 1950s.

Machine extending

It presents a virtual machine (or extended machine) to users that is much easier to access than the underlying physical machine.



The virtual machine presented to the user provides an abstraction of the computer system. This hides the complexity of the hardware from the user; this means that the user need only be concerned with the details of the hardware if they desire.

This is a way of translating the functions needed by a user from the hardware to a presentable and user-friendly medium. As a result, the operating system (OS) acts as an intermediary layer between the user and machine language.

The benefit of this abstraction can be demonstrated when comparing how computations may be processed with and without an operating system (OS).

Without Operating System (OS)	With Operating System (OS)
<p>The instructions written in machine code or assembly language much interface directly with memory hardware. As such, the memory locations to load the two numbers from must be explicitly defined and the memory location to which the result is stored must also be defined.</p> <p>The example below shows a possible assembly code implementation of a computation that is capable of adding two numbers.</p>	<p>The instructions can be written in a high-level language, such as C++.</p>
LDAA \$80 LDAB \$81 ADDB STAA &55	(load number at memory location 80) (load number at memory location 81) (add these two numbers) (store the sum to memory location 55)
	<pre>int a, b, c; a = 1; b = 2; c = a + b;</pre>

Figure 1.9: Adding two numbers

This demonstrates that program development is much more user-friendly with an oper-

ating system (OS). This is because without an operating system (OS) the user must have knowledge of the system hardware; in this case, the necessary memory locations.

In addition, it is possible that the machine code or assembly language written may not work on another computer system as that computer system may have a different architecture or the memory locations may be different. For example, in another computer system:

- memory location 81 may not exist as the memory is smaller; or
- memory location 55 contains important data or instructions that should not be overwritten and therefore the computer system may crash.

By contrast, with an operating system (OS), it is possible to perform computations without interfacing directly with hardware. In the example above, variables (a, b and c) can be used to access data in memory rather than addressing memory locations. The only concern here is that the variable c is able to store the result of a + b.

The operating system (OS) provides a unified environment to users to run their computations in different systems. The operating system (OS) is capable of taking high-level code and translating it in to the machine code that can be executed on a particular computer system.

1.5 Current operating system (OS) trends

Hardware evolution

Due to fast rate of hardware evolution, operating systems (OSs) are more wide-spread than just traditional desktop computers. They can be found on hardware such as:

- mobile devices, such as smartphones and tablets; and
- embedded systems.

Desktop	Mobile	Embedded
<ul style="list-style-type: none">• Windows• macOS• Linux	<ul style="list-style-type: none">• iOS• Android• Symbian OS	<ul style="list-style-type: none">• Windows Embedded

Figure 1.10: Popular operating systems (OSs)

Multiprocessor systems

Definiton

A **multiple processor computer system** makes use of two or more processors and has the ability to allocate tasks between them.

Workstations

A single machine may contain multiple processors and therefore have large computing power.

Distributed and network systems

These computer systems share computing power and peripherals.



The diagram shows an abstraction of a distributed or network system. P1 and P2 are the connected computer systems that both share memory, access to the disk control and, if the system is a network system, the network interface.

However, there is a distinguishable difference between distributed and network systems.

Similarities	Differences
Consist of multiple systems that are interconnected to exchange information.	In distributed systems, users are not aware of the multiplicity of computer systems available.
	In network systems, users explicitly move/share files, submit jobs for processing and other perform other similar tasks.
	In distributed systems, tasks such as moving/sharing files, submitting jobs for processing and other similar tasks are handled automatically by the operating system (OS).

Evaluation

Advantages	Disadvantages
Increase processor throughput due to the use of parallel processing.	A more complex operating system is required in order to be able to interface and manage multiple processor units.
Lower cost than using multiple processors across multiple computer systems because the processors share resources such as the power supply and motherboard.	
Increased reliability because failure of one processor does not affect the other processors and will only slow down the computer system.	

1.6 Operating system (OS) layers

User interfaces

In an operating system (OS), the top layer is the user interface. This is the only layer explicitly visible to the user.

The user interface may be a:

- terminal
 - graphical user interface (GUI)
- text-based command prompt; and/or
 - a visual way of interacting with a computer using items such as windows, icons and menus.

History of the graphical user interface (GUI)

The first company to develop a graphical user interface (GUI) was Xerox PARC. They developed the “Alto” personal computer. It had a bitmapped screen and was the first computer to have a “desktop” screen with a graphical user interface (GUI).

The “Alto” personal computer was not a commercial product. However, several thousand units were manufactured and used at Xerox’s offices and several universities.

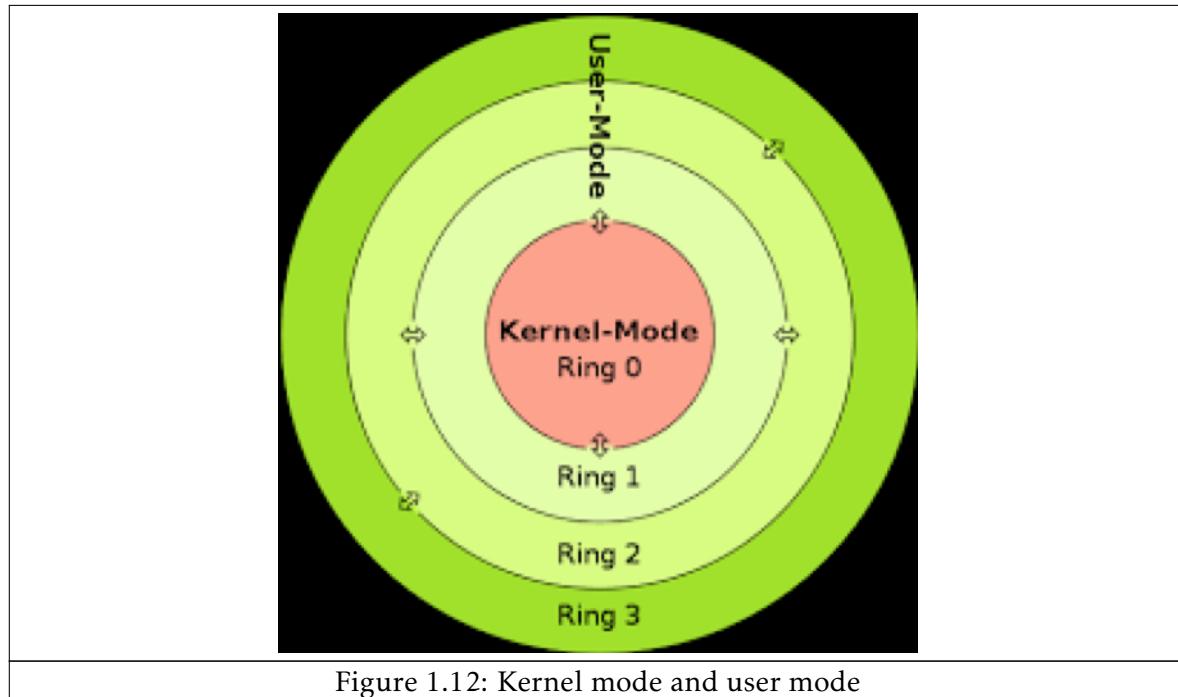
This development was a large influence on the design of personal computers during the late 1970s and early 1980s. Notable examples include:

- Three Rivers PERQ;
- Apple Lisa;
- Apple Macintosh; and
- the first Sun Workstations.

1.7 Kernel mode

Kernel mode vs user mode

Kernel Mode	User Mode
<p>Operating systems (OSs) run in kernel mode.</p> <p>This allows:</p> <ul style="list-style-type: none"> • execution of privileged machine instructions; and • complete access and control of all the hardware. 	<p>Other software runs in user mode.</p> <p>In this mode, instructions that affect control of the machine are forbidden.</p> <p>For example:</p> <ul style="list-style-type: none"> • web browsers; • e-mail software; and • music players.



Ring 0 represents the kernel mode.

Rings 1-3 represent the user mode.

Kernel mode protection

These rings allow separation between the operating system (OS) and user programs for security and protection purposes.

If user mode had unrestricted access to all of the machine instructions:

- a user could inadvertently obtain a virus or write code that is capable of causing damage to the system, and therefore it is necessary to prevent any instructions from directly controlling the machine;
- a program may use resources unfairly, such as holding the CPU or memory, and therefore harm the execution of other programs; and/or
- sensitive machine instructions could be used improperly which may lead to kernel mode errors.

Kernel mode errors are catastrophic.

Kernel Mode	User Mode
<p>A kernel panic represents the operating system (OS) attempting to prevent software causing any harm to the computer system and to recover from a kernel mode error on reboot.</p> <p>An example of a kernel panic is the “Blue Screen of Death” (BSOD) in Microsoft’s Windows.</p>	<p>Application errors where an exception was thrown due to an attempt to execute a privileged instruction, this is one that should only be accessed and executed by the kernel mode, represents the operating system (OS) preventing an application from having unrestricted access to all of the machine instructions.</p>

Figure 1.13: Kernel mode errors

System calls

A **system call** is the programmatic way in which a computer program requests a service from the kernel of the operating system on which it is executed. A system call is a way for programs to interact with the operating system.

Some privileged instructions can be called by a programmer via system calls. Operating systems (OSs) contain system calls for which provide a layer of services for user programs to implement some activities/request services. These are usually sensitive or privileged from the kernel.

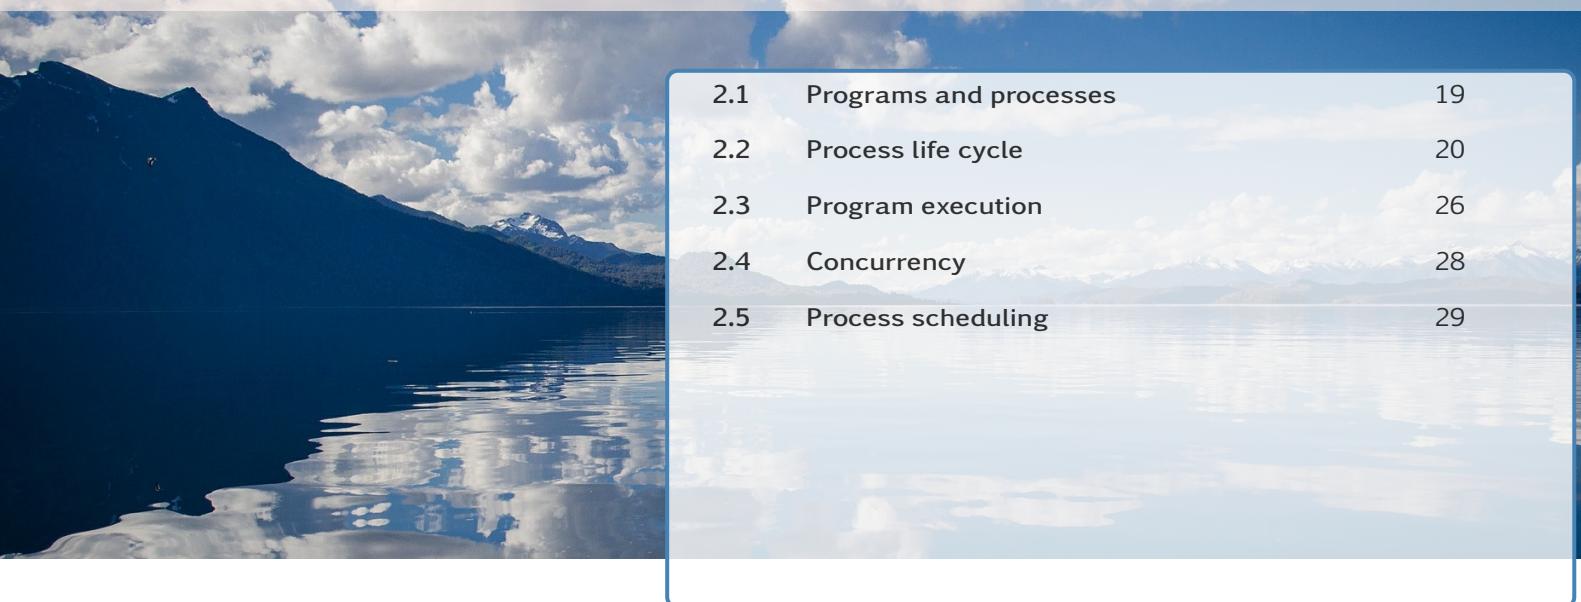
All interactions with the hardware are implemented via system calls. For example, a system call may occur if an application requires interaction with a peripheral such as a printer.

Invoking a system call is similar to calling a general function. However, the difference is that a general function's code is part of program itself, while the system call code is part of the operating system (OS). Different operating systems (OSs) offer different (limited) sets of system calls.

Call	Description
Process Management	
pid = fork()	Create a child process identical to the parent.
exit(status)	Terminate process execution and return status.
File Management	
n = read(fd, buffer, nbytes)	Read data from a file in to a buffer.
n = write(fd, buffer, nbytes)	Write data to a file
Process Management	
seconds = time(&seconds)	Get the elapsed time since Jan 1. 1970

Figure 1.14: Unix system calls

2. Process Management



2.1	Programs and processes	19
2.2	Process life cycle	20
2.3	Program execution	26
2.4	Concurrency	28
2.5	Process scheduling	29

2.1 Programs and processes

Definitions

A **program** is the code written by a programmer.

A **process** (or job/task) shows a program in execution and is a particular instance of a program. These may be shown in a monitoring software, such as Windows Task Manager.

Data are stored values used for the computations by the process.

How it works

A single program may have multiple processes that are currently running.

Each process can share the same code for the program. This is possible as each process uses its own address space, a list of memory locations which the process can read and write. These memory locations contain the program's code instructions and data.

A program is only code however, once it is run, a process is started, and it becomes instructions and data.



2.2 Process life cycle

Process creation

A process can be created by:

- a user – a program may be executed by a user, such as via a double-click using a graphical user interface (GUI) or by typing in a command, and “trigger” the processor to load the program’s executable file containing the program code; or
- another process – an existing process may create another process by spawning/forking – the process that creates a new process is called the parent process while the new process is called the child process. The child process may also spawn a new process forming a tree of processes, as demonstrated in the diagram below.



Process table and process control block

A **process identification number (PID)** is a unique identifier given to a new process when it is created.

A **process control block (PCB)** holds all of the information about a process. It is created when new process is created.

A **process table** stores the process identification numbers (PIDs) for a process and a pointer to the respective process control block (PCB) for that process. This is managed by the operating system (OS).



Figure 2.3: Process table with respective process control blocks

The process descriptor fields in the process control block (PCB) may differ between operating system (OS). An example of possible process descriptor fields is shown by those used by the Minix operating system (OS) are shown below.

Process Management	Memory Management	File Management
Registers	Pointer to text segment	UMASK mask
Program counter	Pointer to data segment	Root directory
Program status word	Pointer to bss segment	Working directory
Stack pointer	Exit status	File descriptors
Process state	Signal status	Effective uid
Time when process started	Process ID	Effective gid
CPU time used	Parent process	System call parameters
Children's CPU time	Process group	Various flag bits
Time of next alarm	Real uid	
Message queue pointers	Effective uid	
Pending signal bits	Real gid	
Process ID	Effective gid	
Various flag bits	Various flag bits	

Figure 2.4: Process descriptor fields in Minix

Although Minix is a fully-featured operating system (OS), it is a small operating system

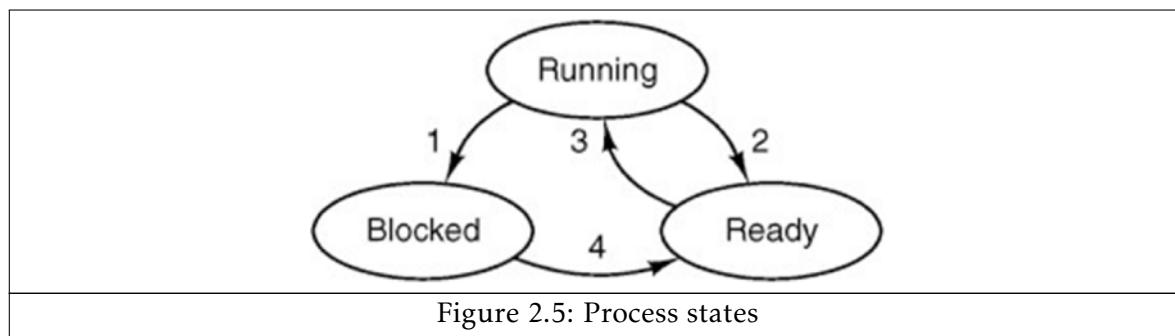
(OS) and therefore the processor descriptor fields are less complex than other operating systems (OSs).

Three-state model

The **three-state model** shows how a process, once initiated, can be in one of three states. The current state of a process is stored in its respective process control block (PCB).

A process, once initiated, can be in one of three main states:

- running – actually using the CPU to perform a task;
- ready – ready to run but waiting for the CPU as it has not yet had time on the CPU has been temporarily stopped to let another process run; or
- ready – unable to run until some external event occurs, such as:
 - waiting for an interrupt, this is a message from the hardware saying that a resource is now available to read from – such as waiting for an input/output (I/O) operation to complete; and
 - waiting for another process to finish accessing a shared resources – for example: a file; memory; or an external peripheral, such as a printer.



In reference to the diagram above, transitions can occur between process states.

Transition	Diagram Number	Explanation
running → blocked	1	Process blocks for input. For example, if the process is waiting for some input from I/O.
running → ready	2	Scheduler picks another process. The process has had opportunity to run and is flagged as no longer currently running, so that another process can run.
ready → running	3	Scheduler picks this process. The next process that is ready to run is set to running to allow access to the CPU.
blocked → ready	4	Input becomes available. The process has received input from I/O or the process sends an interrupt and the interrupt service routine (ISR) is executed, the scheduler is called to transition the process from blocked to ready.

Figure 2.6: Transitions between processor states

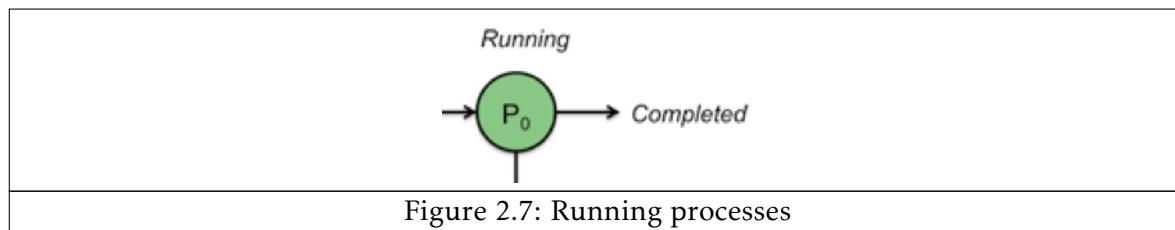
The transition between process states is dependent on the scheduling algorithm used. A clock is used to send a signal to stop the current process, move it from running to ready and then run the scheduler to find out what process should be processed next and the next process will be made to running.

State queues

At any time, a process is in only one state.

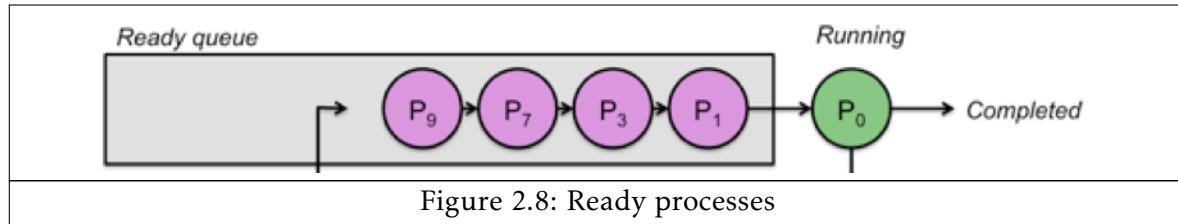
Running processes

At any time, at most one process is in the running state. This is because a single-core processor is only able to process one instruction coming from a single core at a time. If a computer system has a multi-core processor, this rule applies to each individual core on the processor, rather than the processor as a whole, as they are able to complete parallel execution.



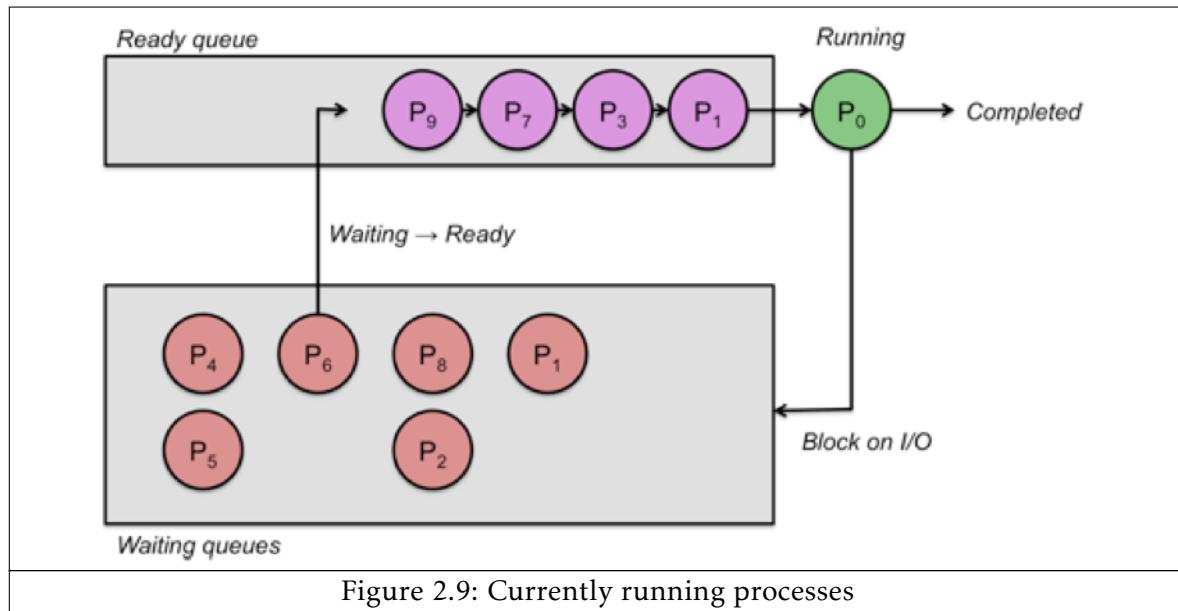
Ready processes

There may be a queue of processes in the ready state.



Blocked processes

There may be several queues of processes in the blocked state, where each queue represents one resource for which processes in that queue are waiting.



Process termination

Process termination is the end of life for a process, this can occur in two general ways.

Voluntary termination

Voluntary termination of a process represents the end of life for a process where its termination was intended by the user or the programmer.

This can be a:

- **normal exit** – the process has done its work; or

- error exit**
 - the process itself handles and “catches” an error – for example, some try-catch code is implemented to check if a condition is met, such as if the definition of a variable is present.



Figure 2.10: Voluntary termination due to an error exit

Involuntary termination

Involuntary termination of a process represents the end of life for a process where its termination is not intended by the user or the programmer.

This can be due to:

- a fatal error**
 - an error is detected by the operating system's (OS's) protected mode – for example, an exception has been thrown due to reference to a non-existent memory location or division by zero; or
- being killed by another process**
 - a process may execute a system call that causes the operating system (OS) to kill another process – this process may have control over the killed process and this may be due to that process being the parent process of the killed process.



Figure 2.11: Voluntary termination due to an error exit

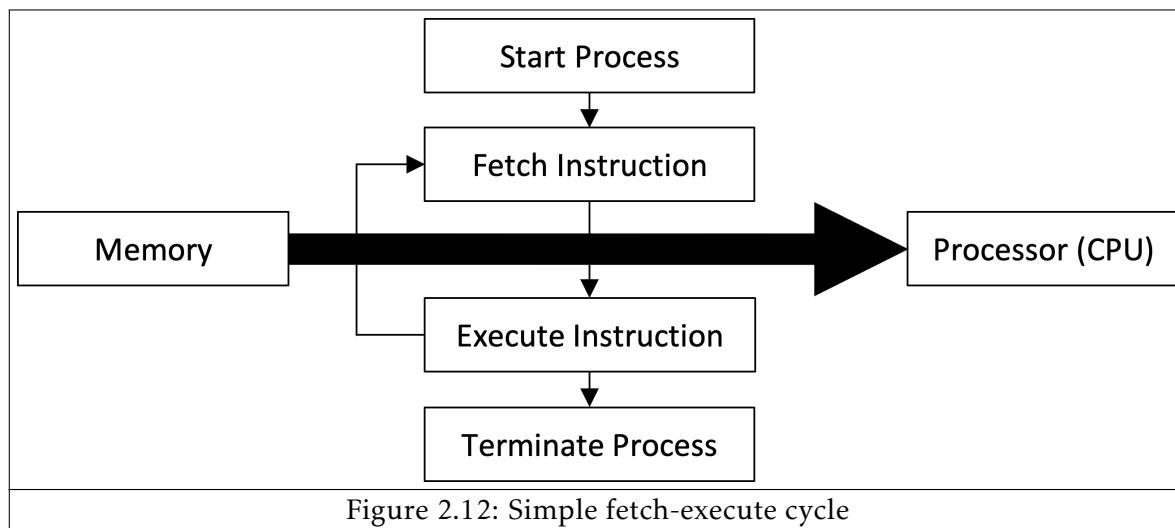
2.3 Program execution

The simple fetch-execute cycle

Definition

The **fetch-execute cycle** is an operational process in which a computer system retrieves a program instruction from its memory, determines what actions the instruction dictates, and carries out those actions.

How it works



Once a process is started, instructions are fetched from memory and executed using the CPU. This process will continue until the process is terminated.

This predictable cycle is only feasible to an extent as some processes may be slow or blocked and others may require immediate attention and cannot wait for the current process to terminate. For example, if a process blocks the processor (CPU) because it is waiting for an event to occur, such as a printer to finish its job, or if a high-priority process is supposed to execute as soon as possible. In these cases, interrupts are required.

Interrupts

Definition

An **interrupt** is a signal sent to the processor indicating that an event caused by hardware or software requires immediate attention.

Types of interrupts

- **Input/output (I/O) interrupt** – Caused by an input/output (I/O) device to signal completion or an error.
- **Timer interrupt** – Caused by a processor timer and is used to alert the operating system (OS) at specific instants.
- **Program interrupt** – Caused by error conditions within user programs or fatal errors.

When do interrupts occur?



Interrupts enable operating systems (OSs) to oversee several programs and input/output (I/O) events simultaneously.

This also means that single-core processors can effectively emulate the way in which multi-core processors deal with multiple instructions at a given time by switching between instructions intelligently. Due to the high clock speeds of modern processors, it is easy to give the illusion that true multi-tasking, where two instructions are being processed at once, is taking place on a single-core processor.

Updated fetch-execute cycle

Given the introduction of interrupts, it is now necessary to update the fetch-execute cycle in order to include the possibility of an interrupt.

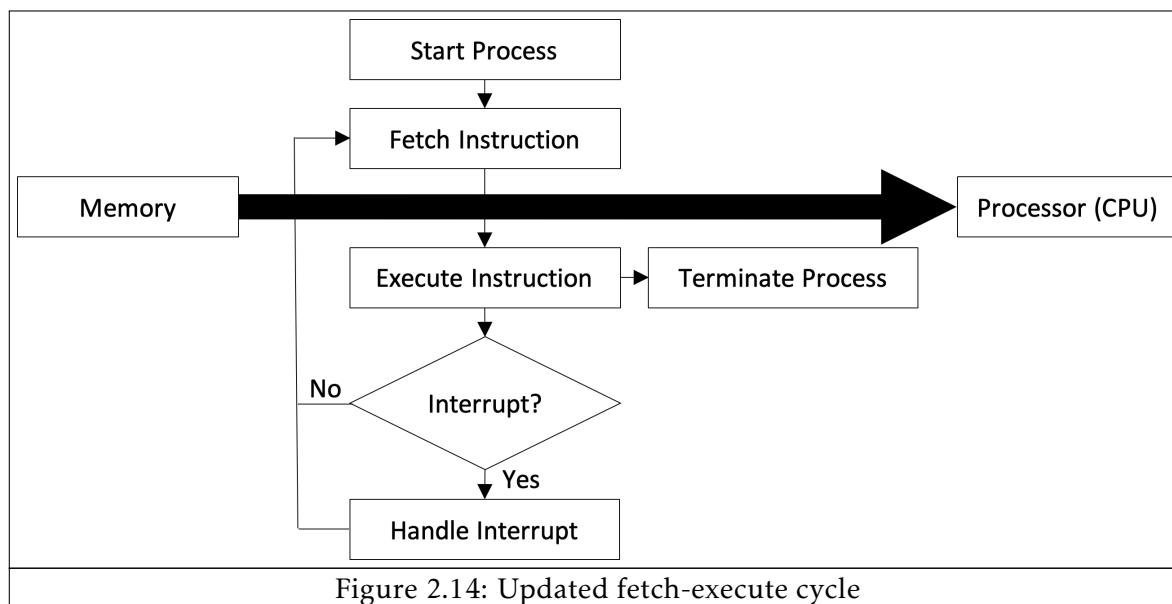


Figure 2.14: Updated fetch-execute cycle

2.4 Concurrency

Definition

Concurrency describes the ability for a program to be decomposed into parts that can run independently of each other. This means that tasks can be executed out of order and the result would still be the same as if they are executed in order.

Why is concurrency required

Concurrency allows the processor (CPU) to run several processes. An example of this is shown by an interrupt occurring due to an input/output (I/O) event occurring (page 24).

Interleaving

Concurrency is able to achieve multitasking, that does not require parallel execution, by performing interleaved execution.



2.5 Process scheduling

The scheduler

Definition

A **scheduler** uses a scheduling algorithm to determine how to share processor time.

How it works

After an input/output (I/O) system call or interrupt handling, control is passed to the scheduler to decide which process to execute next.



The scheduler checks if the current process is still the most suitable to run at this moment in time. If it is control is returned to the process, otherwise:

- the state of the current process is saved in the process control block (PCB);
- the state of the most suitable process is retrieved from the process control block (PCB); and
- control is transferred to the newly selected process at the point indicated by the restored program counter (PC).

The action of storing the state of the current process and activating another process is called a context switch.

Context switching must be minimised to reduce overheads created by copying the state of processes and the time taken to perform the switch. However, it is still regarded as important to context switch when appropriate to avoid longer waiting times in the event of a blocked process.

Scheduling

Definition

Scheduling is the act of determining the optimal sequence and timing of assigning execution to processes.

Scheduling criteria

Different scheduling criteria may be selected depending on the use case for a given computer system.

- **CPU utilisation** – Aims to keep the CPU as busy as possible.
- **Efficiency** – Aims to maximise system throughput.
- **Fairness** – Aims to be fair to all running processes or to all users on a multi-user operating system (OS).

This means that different policies and algorithms for scheduling will exist to match these criteria.

Scheduling policies

A **non-preemptive scheduling policy** is one that allows processes to run until complete or incurring an input/output (I/O) wait. These scheduling policies can be described as passive.

A **preemptive scheduling policy** is one that allows processes to be interrupted and replaced by other processes, generally through timer interrupts.

Scheduling algorithms

Definitions

Arrival time is the instant at which a process is first created.

Service time is the time that it takes for a process to complete if it is in continuous execution.

The **waiting time** for a process is the sum of time spent in the ready queue during the life of the process. This does not include time that the process is blocked or waiting for input/output (I/O).

In order for a scheduling algorithm to be deemed as fair to all processes, the ratio between waiting time and run time should be about the same for each process.

Case study

Process	Arrival Time	Service Time
A	0	3
B	2	6
C	4	4
D	6	5
E	8	2

Figure 2.17: Case study

The case study shows different processes (A-E) that all have different arrival times and different service times.

The following examples of scheduling algorithms will refer to this case study to demonstrate their function.

First come, first served (FCFS) / First in, first out (FIFO) – Non-preemptive

In this algorithm, the first process to arrive is assigned to the processor (CPU) until it is finished. Meanwhile, any other processes that come along are queued up waiting to be processed.



Advantages	Disadvantages
Simple to implement.	Does not consider the priority of a process and therefore the important processes may not be completed quickly.

	Prevents other processes from starting while another is in progress and therefore, if processes are of varying sizes, there may be inefficiencies since a single process may take a long time to complete therefore leaving the user waiting before they can perform any other actions.
Favours long processes as all processes are given the opportunity to run until completion and therefore, some shorter processes may not be able to start processing until the longer processes are finished.	

Shortest job first (SJF) – Non-preemptive

In this algorithm, the process with the shortest estimated run time is assigned to the processor (CPU) until it is finished. Meanwhile, any other processes that come along are queued up waiting to be processed.

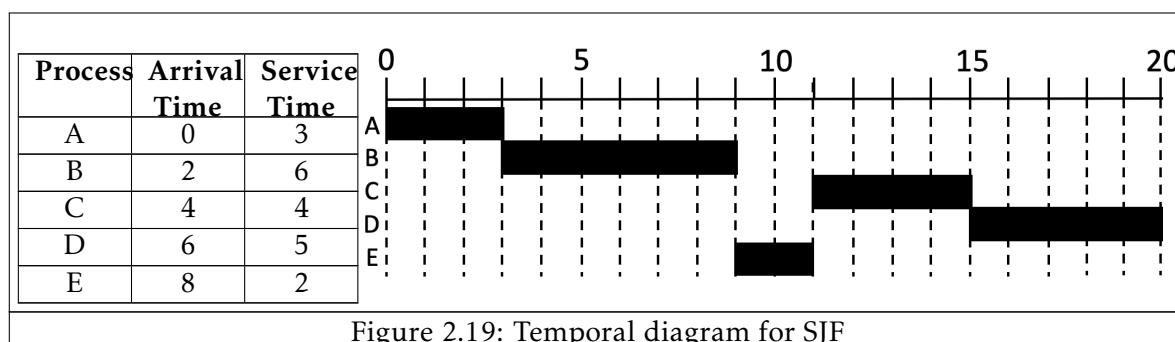


Figure 2.19: Temporal diagram for SJF

Advantages	Disadvantages
Simple to implement.	Does not consider the priority of a process and therefore the important processes may not be completed quickly.
Shorter processes are processed quickly because they take precedence.	Relies on an estimation of how long a process will take which could be incorrect.
Minimises the average time taken to complete a process because the shortest processes take precedence.	Prevents other processes from starting while another is in progress and therefore, if processes are of varying sizes, there may be inefficiencies since a single process may take a long time to complete therefore leaving the user waiting before they can perform any other actions.
Favours long processes as the processes with the shortest estimated time are given the opportunity to run until completion and therefore, some longer processes may not be able to start processing until there are no more shorter processes currently running or ready to be processed.	

Shortest remaining time (SRT) – Preemptive

In this algorithm, the process with the shortest estimated remaining time is assigned to the processor (CPU). If a process becomes ready that has a shorter remaining time, it will be pre-empted to allow the new process to start and the scheduler is switched to that new process.



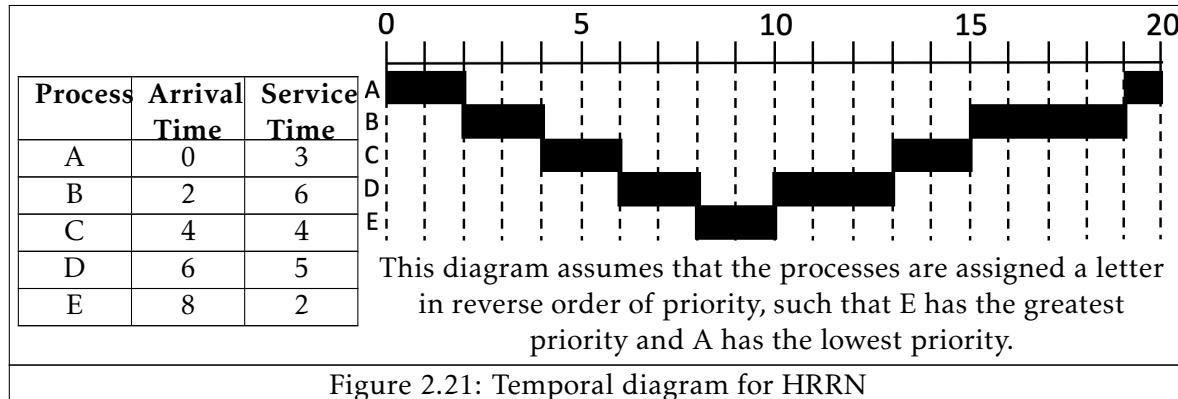
Advantages	Disadvantages
High throughput because the number of processes completed is high due to the shortest processes taking precedence.	Does not consider the priority of a process and therefore the important processes may not be completed quickly.
Shorter processes are processed quickly because they take precedence.	Relies on an estimation of how long a process will take which could be incorrect.
Favours long processes as the processes with the shortest estimated time are given the opportunity to run until a process with a shorter estimated time is ready and therefore, some longer processes may not be able to start processing until there are no more shorter processes currently running or ready to be processed.	

Highest response ratio next (HRRN) – Preemptive

In this algorithm, the process with the highest priority is assigned to the processor (CPU). If a process becomes ready that has a higher priority, it will be pre-empted to allow the new process to start and the scheduler is switched to that new process.

The priority of a process may be based on memory, time and/or any other resource requirement such as:

$$\frac{\text{waiting time} + \text{runtime}}{\text{runtime}}$$



Advantages	Disadvantages
High throughput because the number of processes completed is high due to the shortest processes taking precedence.	Does not consider the priority of a process and therefore the important processes may not be completed quickly.
Shorter processes are processed quickly because they take precedence.	Relies on an estimation of how long a process will take which could be incorrect.
	Can be inefficient if a large process is in progress and shorter processes are being added to the queue because they will take precedence.

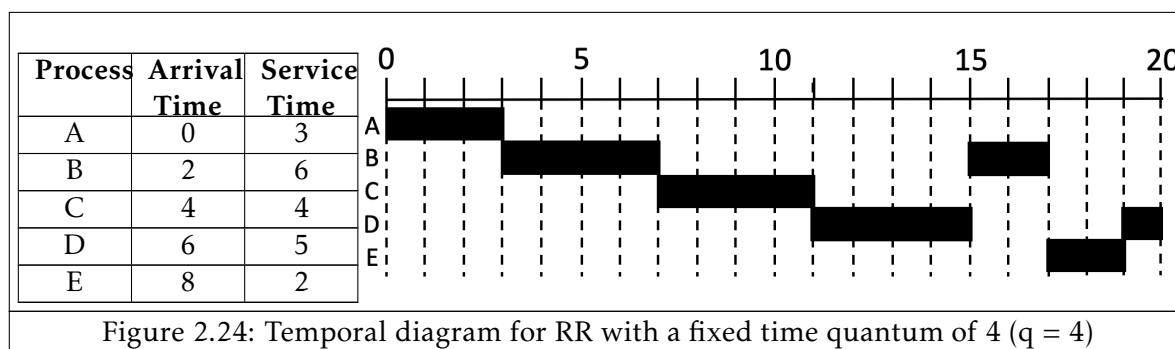
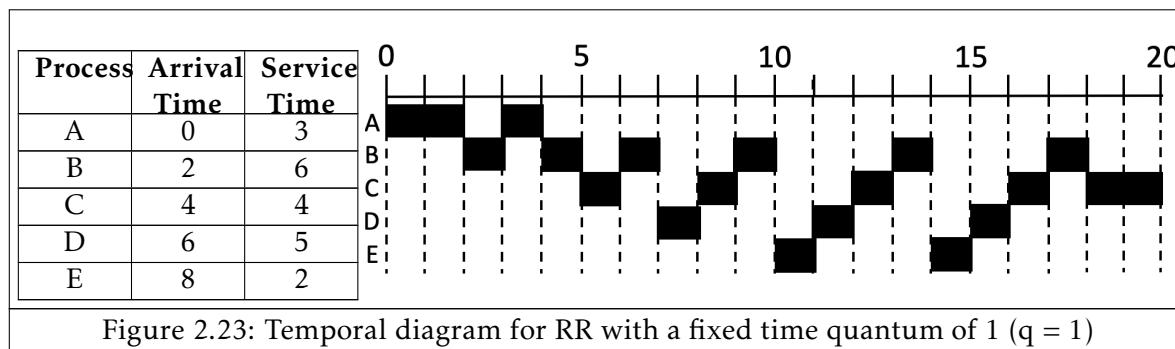
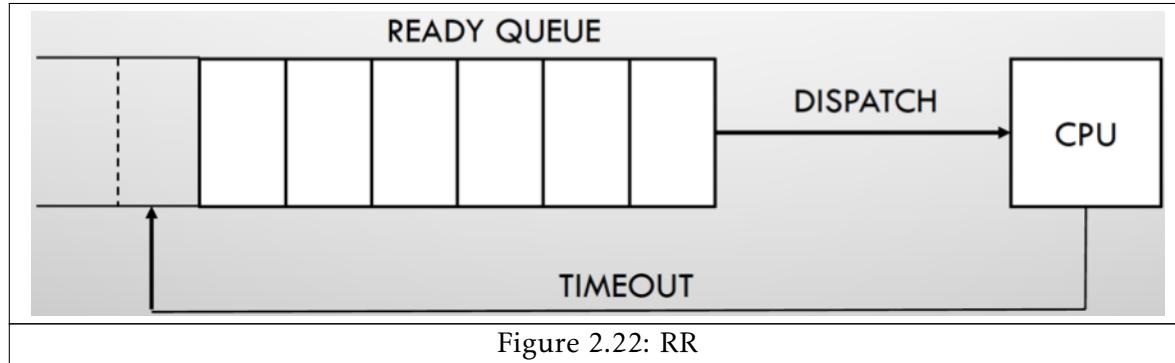
Round robin (RR) – Preemptive

In this algorithm, each process is dispatched to the processor (CPU) on a “first in, first out” (FIFO) basis with a fixed time quantum.

Each time quantum is typically 10-20ms. Modern processor (CPU) clock frequencies are typically greater than 2GHz, which implies clock periods of 5×10^{-7} ms. This shows that the given time quantum is relatively large compared to a typical processor’s (CPU’s) clock period.

If a process experiences a timeout, this will mean that it has run over its fixed time quantum. In which case, the process will be interrupted and returned to the back of the queue.

A system designer may wish to choose a time quantum that is most appropriate for a given system. This may be done by measuring the average service time and waiting time for the processes that will be running on the system and design the fixed time quantum around these figures. It may be that a system designer wishes to minimise the number of processes that are interrupted by another process.



Advantages	Disadvantages
Simple to implement.	Heavy overhead due to continuous context switches.
Suitable for some types of computer systems , such as those which will be running processes of similar priority and size.	Does not consider the priority of a process and therefore the important processes may not be completed quickly.
	Does not consider the size of a process and therefore, if processes are of varying sizes, there may be inefficiencies since a single process may take a long time to complete thus leaving the user waiting before they can perform any other actions.

Multilevel queueing

Definition

Multilevel queueing is a queue with a predefined number of levels which consist of several independent queues.

How it works

Multilevel queueing makes use of other existing scheduling algorithms.



Figure 2.25: Multilevel queueing

Each queue has a different priority; the top queue has the highest priority and the bottom queue has the lowest priority. Processes are able to move between queues if their priority changes.

There is some form of inter-queue scheduling policy that governs the assignment of processes from each queue to the processor (CPU).

The queues in a multilevel queueing system may differ between operating system (OS). For example, the scheduler used by the Minix operating system (OS) uses multi-level queueing and implements 16 queues.

Advantages	Disadvantages
Allows scheduling optimisation as a system designer may be able to leverage the advantages of a range of different scheduling algorithms.	
Maintains common processes as it is possible to split processes into different queues depending on their nature. For example, input/output processes could be in one queue while processor (CPU) processes are in another queue.	

Helps to prevent bottlenecks because input/output (I/O) devices are slower than the processor speed and therefore maximising processes involving input/output (I/O) devices ensures that these devices are continuously busy.	
--	--

3. Threads and Concurrency



3.1	What are threads?	39
3.2	Sequential and concurrent programming	43
3.3	Sequential execution	45
3.4	Introduction to concurrent execution	47
3.5	Interleaving	48
3.6	User and kernel threads	52
3.7	Multi-threading models	54
3.8	Evaluation of concurrent programming	57

3.1 What are threads?

Definition

A **thread** is an independent path/sequence of execution within a process and can be managed independently by a scheduler. A process may contain many threads.

Multi-threading

How it works

As mentioned in the previous section, a process (or job/task) shows a program in execution and is a particular instance of a program. By default, these processes are run by means of “single execution thread”.

However, different parts of the same process could be parallelised in order to allow multi-threading.

Multi-threading provides a way of improving application performance and therefore improving the efficiency and/or usability of a computer system.



Figure 3.1: Single-threaded vs multi-threaded

Example

A screenshot of Microsoft Word's 'Document1 - Microsoft Word' window. The main content area displays two paragraphs of text. Below the text, there is a 'Word Processor Checklist' section with links to 'Jump to a topic in this article' and 'History of Word Processing' and 'Standard Features of Word Processors'. The top of the window shows the Microsoft Word ribbon with tabs like File, Home, Insert, Page Layout, References, Mailings, Review, and View. The Home tab is selected, showing font and paragraph formatting tools. To the right of the checklist, there is explanatory text: 'In this example, a thread could be assigned to each of the following tasks:' followed by a bulleted list of tasks, and 'This is possible as all of these tasks can be run in parallel as they do not require information from one another.'

Figure 3.2: Example of multi-threading in Microsoft Word

Threads vs processes

Comparison

Threads and processes share some similarities however, there are also some distinct differences.



Figure 3.3: Single-threaded process

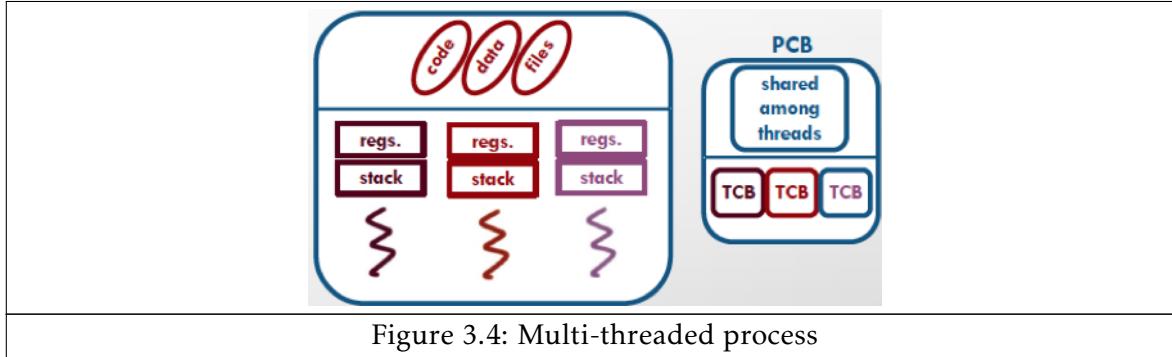


Figure 3.4: Multi-threaded process

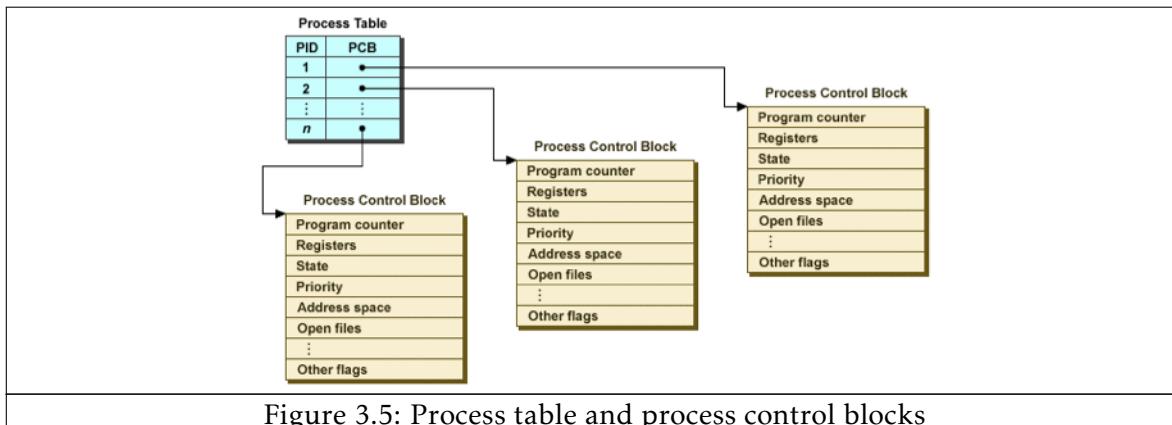


Figure 3.5: Process table and process control blocks

	Threads	Processes
Similarities	Sequential flow of control with start and end.	Sequential flow of control with start and end.
	At any time, a thread has a single point of execution.	At any time, a process has a single point of execution.
	Has its own execution context, stack (history) and program counter stored in a thread control block (TCB).	Has its own execution context, stack (history) and program counter stored in a process control block (PCB).
	Follows the three-state model in which the thread can be running, blocked or ready.	Follows the three-state model in which the process can be running, blocked or ready.
	Context switching can happen for threads.	Context switching can happen for processes.
	A thread can spawn another thread.	A process can spawn another process.
	A thread is often called a lightweight process.	
Differences	A thread cannot exist on its own, instead it exists within a process.	A process does not require a parent entity.
	Usually created and/or controlled by a process.	A process is not typically created and/or controlled by another process.
	Threads can share process properties, including memory and open files.	Processes cannot share process properties with other processes.
	Inexpensive creation and context switching as does not require separate address space.	Expensive creation and context switching as requires separate address space.
	When running multiple threads concurrently, they share an address space.	When running multiple processes concurrently, they are resources, such as memory, disk and printers.

Figure 3.6: Threads vs processes

Properties

Per process items	Per thread items
Address space	Program counter
Global variables	Registers
Open files	Stack
Child processes	State
Pending alarms	
Signals and signal handlers	
Accounting information	

Figure 3.7: Threads vs processes

Process properties are shared between threads. Thread properties are local and private to each thread.

3.2 Sequential and concurrent programming

Sequential programming

Definition

Sequential programming is the traditional activity of constructing a computer program using a sequential programming language.

How it works

This involves a programming methodology that assumes statements are executed in order/sequence.

Programs written using sequential programming are assumed to execute on a single-CPU system and have a single thread of control.



Evaluation

Advantages	Disadvantages
No additional support required from the programming language.	Lower processor throughput than concurrent programming as it cannot benefit from multitasking or concurrent processing.
No additional support required from the operating system as most old-school operating systems were generally single-threaded and therefore later generations of operating systems typically inherit this functionality.	Multiple computer systems that each have their own CPU may yield a higher cost than multi-CPU systems as more will be spent on resources, such as the power supply and motherboard, that could be otherwise shared in a multi-CPU system.
	Lower reliability than a multi-CPU system because, in the case of the failure of the processor, there is no redundancy.

Concurrent programming

Definition

Concurrent programming is the activity of constructing a computer program that takes advantage of concurrency allowed by the use of multiple threads of control.

How it works

Multiple threads of control allow a given process to perform multiple computations in parallel and to control simultaneous external activities.

The program may be run on both:

- a single-CPU – where the computer program will take advantage of multitasking; system and
- a multi-CPU – where the computer program will take advantage of true system parallelism.

Evaluation

Advantages	Disadvantages
Increase processor throughput due to the use of multitasking in a single-CPU system or parallel processing in a multi-CPU system.	Requires support from the programming language as it must implement techniques to deal with multitasking on a single-CPU system and/or parallel processing on a multi-CPU system.
A multi-CPU system generally yields a lower cost than using multiple CPUs across multiple computer systems because the processors share resources such as the power supply and motherboard.	Requires support from the operating system as it must support multi-threading in order to allow multitasking on a single-CPU system and/or interface and manage multiple CPU units on a multi-CPU system.
Increased reliability in a multi-CPU system because failure of one processor does not affect the other processors, instead the computer system may experience lower performance until fixed.	

3.3 Sequential execution

Definition

Sequential execution is where the execution of threads in a sequential program is executed in sequence/order with no overlapping.

Order and precedence

Explanation

In sequential execution, there is only one possible sequence of execution. This is because a sequential program gives the system strict instructions on the order of executing the state-

ments in the program.

Importance

For example, a simple hypothetical program could be:

P;

Q;

R;

This tells the computer system that the statements must be executed in the order they are written, such that:

- P must precede Q; and
- Q must precede R.

High level

The importance of the order of precedence can be highlighted by demonstrating this idea in a high-level programming language.

Given the following program written in a high-level language:

```
x = 1;  
y = x + 1;  
x = y + 2;
```

it is possible to see that the final values of x and y depend on the order of execution of the statements.

System level

Given the following program written in a high-level language:

```
x = 1; P  
y = x + 1; Q  
x = y + 2; R
```

where each statement is assigned a letter respectively, each statement may be compiled into several machine instructions.

Statement **P** is treated as a single machine instruction:

- **P1**: store 1 at the memory address of x.

Statement **Q** is broken into three machine instructions:

- **Q1**: load the value of x into a CPU register;

- **Q2:** increment the value in this register by 1; and
- **Q3:** store the value in this register at the memory address of y.

Statement **R** is broken in to three machine instructions:

- **R1:** load the value add of y in to a CPU register;
- **R2:** increment the value in this register by 2; and
- **R3:** store the result at the memory address of x.

The nature of sequential execution

The execution of statements **P**, **Q** and **R** at the program level (or high-level) as

P → **Q** → **R**

implies that the execution at the system level is as follows

P1 → **Q1** → **Q2** → **Q3** → **R1** → **R2** → **R3**,

given that **P** is compiled to a single machine instruction, whilst **Q** and **R** are compiled to three machine instructions – as seen on page 40.

Sequential execution has the following assumptions:

- total ordering
 - there is single-threaded computation, and therefore no overlap in the execution of the statements;
- deterministic
 - the same input will always result in the same output; and
- sequence
 - users will specify a strict sequence of steps required in order to achieve a desired goal.

However, this does not apply in many practical applications, for which a sequence of steps is not required.

3.4 Introduction to concurrent execution

Definition

Concurrent execution is where the execution of threads in a concurrent program is occurring asynchronously, meaning that the order in which tasks are executed is not predetermined.

The squares example

In this hypothetical example, a person desires to have a list with the results of all of the squares (2^n) from 1 to 100000.

A group of 100000 people are split in to heavily uneven teams and assigned the same task to complete all of the calculations in order to achieve the desired result.

It is given that each calculation takes n amount of time.

Team 1		Team 2	
Number of members	1	Number of members	99999
Strategy	One person should complete all of the calculations.	Strategy	Each member is assigned a number between 1 to 100000. Each member should calculate the respective square for the number they are assigned.
Time taken	$100000n$	Time taken	n

This shows that Team 2 was $100000x$ faster than Team 1. This was because it was possible to decompose the larger task in to smaller sub-tasks and assign each of those tasks to a separate resource, which in this case is one person.

This example forms that basic concept of concurrent execution.

The nature of concurrent execution

Concurrent execution dismisses many of the assumptions required for sequential execution (page 41).

Calculations may be carried out without total ordering. As a result, calculations may be carried out in parallel and overlapping is therefore allowed.

In the example above, each individual person in team 2 carried out their operations in sequence.

In the example above, the operations in the whole computation can be viewed as being in partial order. However, the order does not matter here because there is no dependency between the calculations. This is because the output from any given calculation is not required as an input to any other given calculation.

However, in general, concurrent execution is non-deterministic, and therefore the same input generally means different output due to ordering. This is because there are many cases where the order of operation does matter.

3.5 Interleaving

Why is interleaving required?

Concurrent execution on a computer system with a multi-core CPU or multiple CPUs can make use of parallel processing in order to run threads asynchronously. However, this is not possible on a computer system with a single-CPU that consists of only one core.

As a result, interleaving is used in order to switch execution between threads.

It is important to note that the operations within each thread are strictly ordered, but the interleaving of the operations are not ordered and are interleaved in an unpredictable order.

Calculating interleavings

Formula

It is possible to calculate the number of interleavings given the formula:

$$\text{number of interleavings} = \frac{t_1 + t_2 + \dots + t_n}{t_1! t_2! \dots t_n!}$$

where

- t_n represents the number of statements/operations in each thread.

For example, in a concurrent program that has two threads, the formula may be adjusted to:

$$\text{number of interleavings} = \frac{(n+m)!}{n!m!}$$

where

- n represents the number of statements/operations in the first thread (t_1); and
- m represents the number of statements/operations in the second thread (t_2).

Example

A high-level concurrent program may spawn new threads.



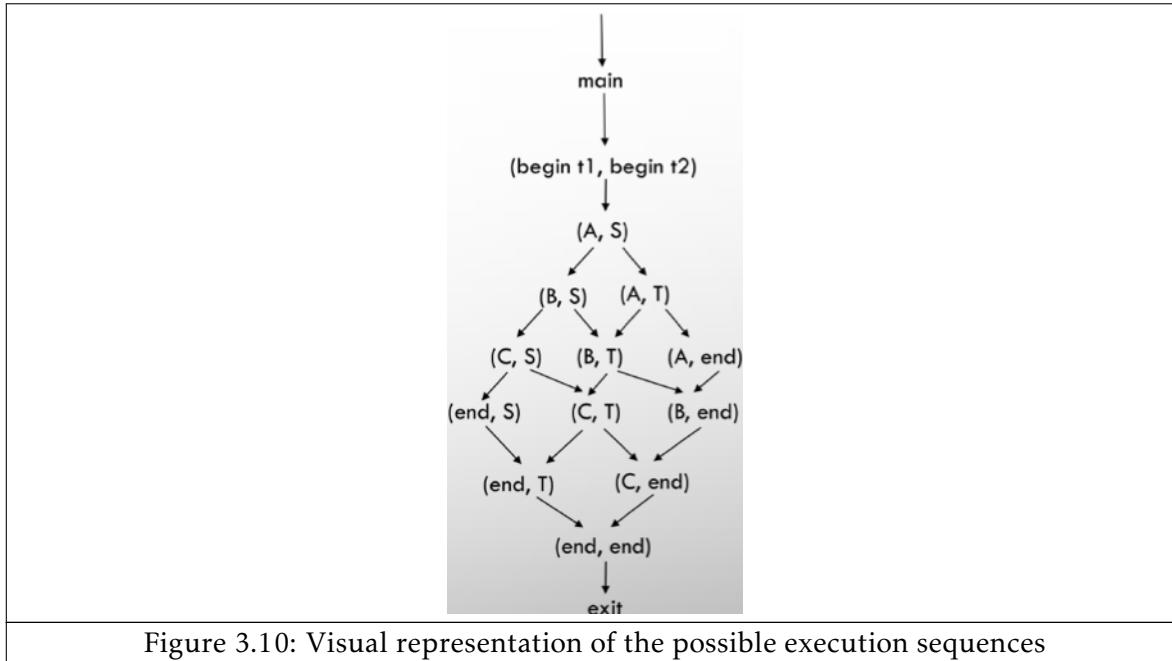
In this example, the concurrent program spawns the threads t_1 and t_2 where:

- t_1 has three statements – A, B and C; and
- t_2 has two statements – S and T.

There are two threads and therefore it is possible to use the adjusted formula to calculate the number of interleavings in this concurrent program:

$$\begin{aligned} \text{number of interleavings} &= \frac{(n+m)!}{n!m!} \\ \text{number of interleavings} &= \frac{(3+2)!}{3! \times 2!} \\ \text{number of interleavings} &= \frac{6!}{3! \times 2!} \\ \text{number of interleavings} &= \frac{6 \times 5 \times 4 \times 3 \times 2 \times 1}{(3 \times 2 \times 1) \times (2 \times 1)} \\ \text{number of interleavings} &= \frac{120}{6 \times 2} \\ \text{number of interleavings} &= \frac{120}{12} \\ \text{number of interleavings} &= 10 \end{aligned}$$

There are 10 possible interleavings, thus yielding 10 possible different execution sequences.



A run of the program corresponds to an interleaving sequence. Each interleaving sequence determines a unique sequence of executing the statements. Repeated runs with the same input will likely trace different interleavings.

Growth of interleavings

The number of interleavings grows extremely quickly given an increase in:

- the number of threads in the concurrent program; or
- the number of statements/operations in one or more of the concurrent program's threads.

This can be demonstrated by increasing the number of operations in the previous example:

- t_1 now has four statements – A, B, C and D; and
- t_2 now has five statements – S, T, U, V and W.

and therefore:

$$\begin{aligned}
 \text{number of interleavings} &= \frac{(n+m)!}{n!m!} \\
 \text{number of interleavings} &= \frac{(4+5)!}{4! \times 5!} \\
 \text{number of interleavings} &= \frac{9!}{4! \times 5!} \\
 \text{number of interleavings} &= \frac{9 \times 8 \times 7 \times 6 \times 5 \times 4 \times 3 \times 2 \times 1}{(4 \times 3 \times 2 \times 1) \times (5 \times 4 \times 3 \times 2 \times 1)} \\
 \text{number of interleavings} &= \frac{362880}{24 \times 120} \\
 \text{number of interleavings} &= \frac{362880}{2880} \\
 \text{number of interleavings} &= 126
 \end{aligned}$$

3.6 User and kernel threads

User threads

User threads are created and managed by a user level library, typically without the knowledge of the kernel.



The diagram shows that:

- all of the threads for a given process is present within the user space; and
- the thread table is present within the process.

User threads are:

- fast to create and manage; and
- portable to any operating system (OS).

If one user thread is blocked, all other threads in the same process are also blocked. For example, in a word processor application, a thread that handles a printing event would block all other threads and therefore prevent the user from interacting with other aspects of the application.

Multi-threaded applications cannot take advantage of parallel execution on computer systems with a multi-core CPU or computer systems with multiple CPUs.

Kernel threads

Kernel threads are directly managed and supported by the operating system's (OS's) kernel.



The diagram shows that:

- all of the threads for a given process is present within the user space; and
- the thread table is present within the kernel space, rather than the process itself or the user space.

Kernel threads are:

- slower to create and manage than user threads; and
- specific to the operating system (OS).

If one user thread is blocked, all other threads in the same process are scheduled and not

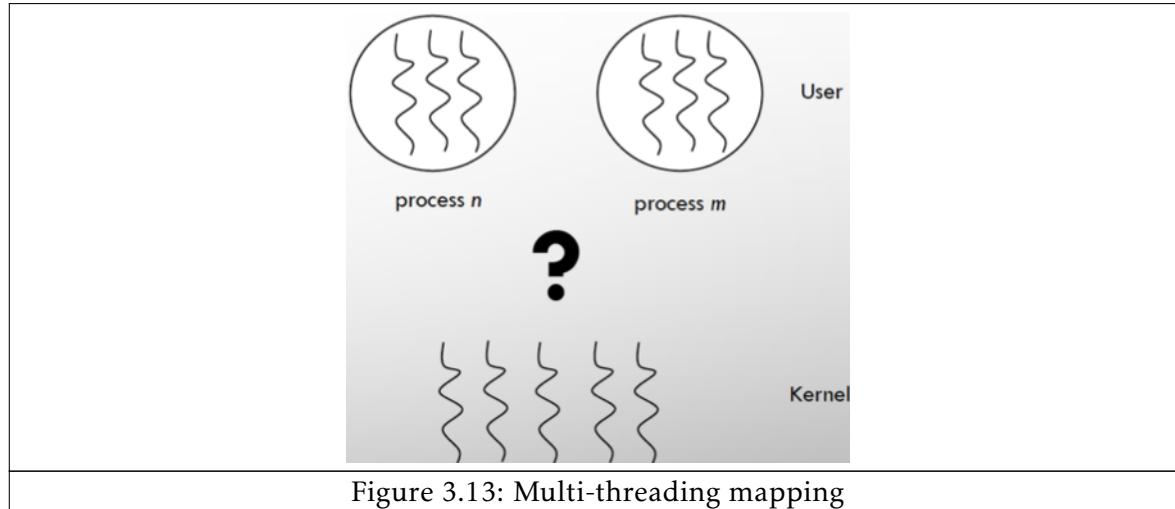
blocked. For example, in a word processor application, a thread that handles a printing event would no longer block all other threads and therefore would allow the user to interacting with other aspects of the application whilst the printing event occurs.

Can take advantage of parallel execution on computer systems with a multi-core CPU or computer systems with multiple CPUs.

3.7 Multi-threading models

Why is multi-threading mapping required?

The kernel is generally not aware of the user threads present in a process. Therefore, a thread library must map user threads to kernel threads.



The diagram shows that there must be some relationship between the user threads and the kernel threads. This relationship may be defined using different mappings, including:

- many-to-one;
- one-to-one; and
- many-to-many.

Many-to-one mapping

How it works

All user threads from each process map to one kernel thread.



Evaluation

Advantages	Disadvantages
Portable as there are few system dependencies.	No parallel execution of threads.
	No concurrency as all threads in a process are blocked if another thread is blocked, for example if the thread is waiting for an input/output (I/O) interrupt.

One-to-one mapping

How it works

Each user thread maps to a single kernel thread.



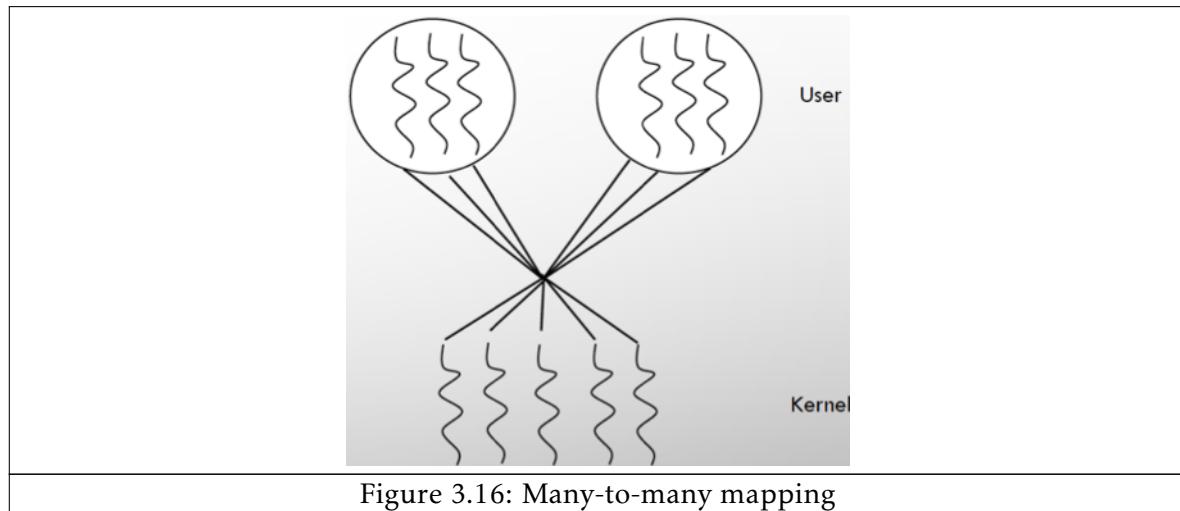
Evaluation

Advantages	Disadvantages
Concurrency as all threads in a process are not blocked if any given thread becomes blocked.	Slow as there is management overhead because the kernel is involved for every user thread.
Performance as it can take advantage of multiple CPUs.	Restricted as there is typically a limit on the number of threads.
	Creating user threads requires the corresponding kernel support.

Many-to-many mapping

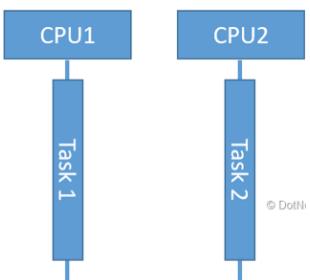
How it works

Many user threads multiplex to an equal or smaller number of kernel threads.



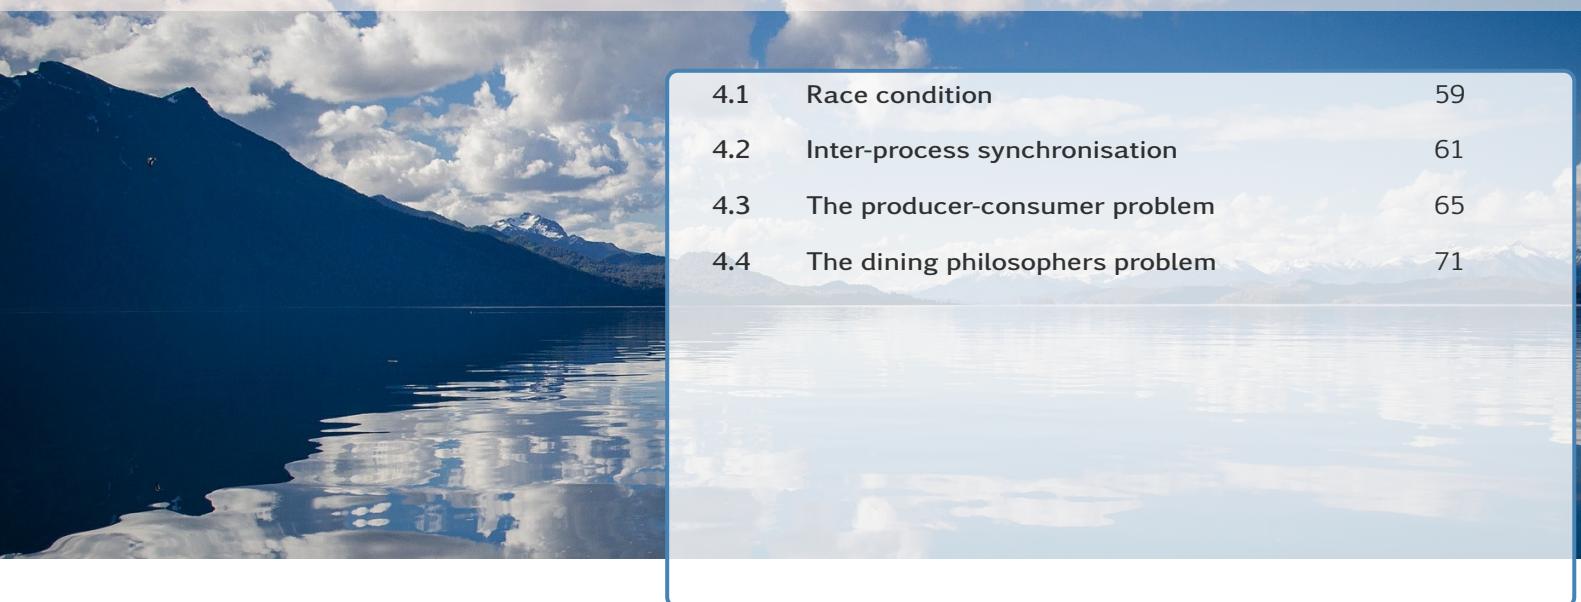
Advantages	Disadvantages
Performance as it can take advantage of multiple CPUs.	Complexity and therefore implementation difficulties.
Flexible as there is no limit on the number of threads.	

3.8 Evaluation of concurrent programming

Advantages	Disadvantages
 <p>Parallelism. It improves the efficiency of program execution in computer systems with multiple CPUs by allowing tasks/operations to be split up and executed independently on each CPU.</p>	<p>Debugging complexity as concurrent programs are non-deterministic and therefore it can be difficult to trace a problem/bug in the code as the same input will generally not result in the same output.</p>
 <p>Multi-tasking. It improves the utilisation of the CPU in a computer system that only has a single CPU. This allows multiple tasks/operations to run alongside each other and appear to be processed simultaneously.</p>	<p>No protection between threads.</p>
<p>Increases application responsiveness, for example, in a word processor application one thread could be responsible for responding to user input/output (I/O) while other threads perform tasks in the background.</p>	<p>Concurrent processes must interact with each other in order to share resources or exchange data.</p>
<p>Suited to some applications as there are some practical applications that are non-deterministic and concurrent as the order of program operations is determined by other external events. This is useful for applications that need to handle multiple events.</p>	<p>Synchronisation must be promoted in order to determine when, how and with what language abstractions computation events can be synchronised in order to eliminate unacceptable outputs.</p>
	<p>Distribution must be taken care of in order to consider how threads can be distributed among a number of CPUs and how one thread is able to interact with another thread on a different CPU.</p>

	<p>Error-prone. Examples of major concurrent programming errors include:</p> <ul style="list-style-type: none">• Therac-25 - A computerised radiation therapy machine whose errors contributed to accidents causing deaths and serious injuries.• Mars Rover Pathfinder – Problems with interaction between concurrent tasks caused periodic software resets, thus reducing availability for exploration.
--	---

4. Synchronisation and Mutual Exclusion



4.1	Race condition	59
4.2	Inter-process synchronisation	61
4.3	The producer-consumer problem	65
4.4	The dining philosophers problem	71

4.1 Race condition

Definition

A **race condition** describes the competition for resources in a critical section caused by interleaving/thread interference.

Why do race conditions happen?

Race conditions occur due to interleaving/thread interference.

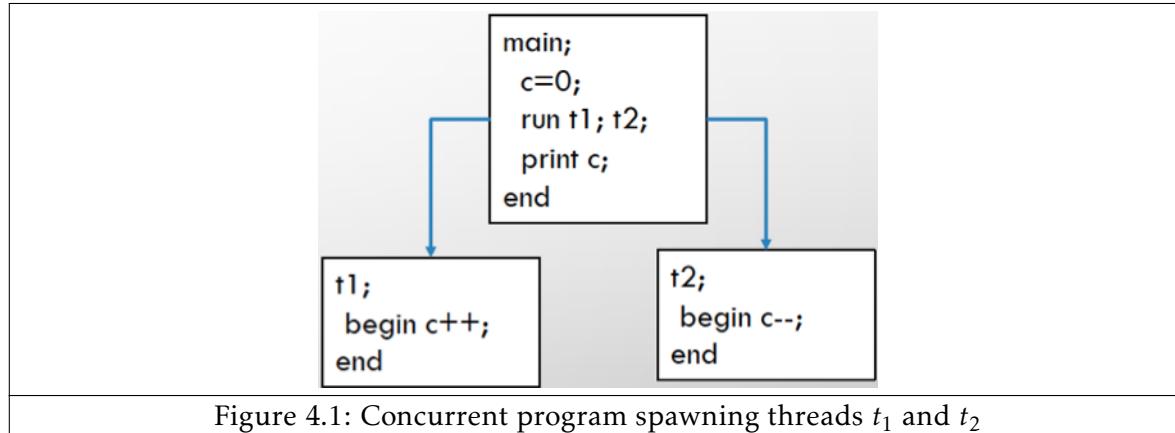
Interleaving/thread interference describes an undesired outcome resulting from non-deterministic, concurrent usage of shared resources.

This happens because, in general, concurrent execution is non-deterministic, and therefore the same input generally means different output due to ordering. This is because there are many cases where the order of operation does matter.

Examples

Racing for memory access

A race condition may occur when two threads attempt to access the same location memory, such as registers or RAM, at the same time.



In this example, the two threads t_1 and t_2 manipulate the same variable where:

- t_1 increments the variable c ; and
- t_2 decrements the variable c .

As seen before (page 40), each statement may be compiled in to several machine instructions.

The increment ($c++$) instruction is broken in to three machine instructions:

- retrieve c ;
- increment retrieved value; and
- store result in c .

The decrement ($c--$) instruction is also broken in to three machine instructions:

- retrieve c ;
- decrement retrieved value; and
- store result in c .

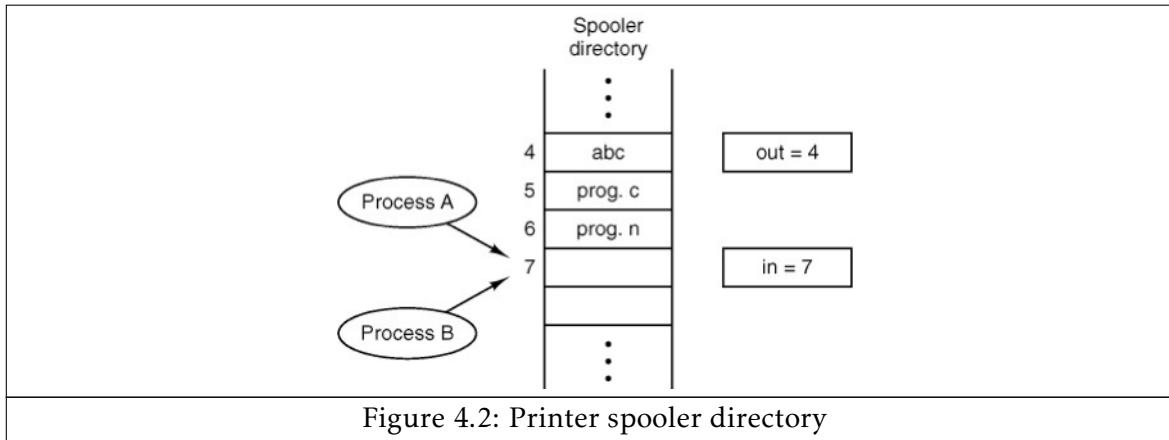
As a result, one interleaving possibility is as follows:

- t_1 : retrieve c ;
- t_2 : retrieve c ;
- t_1 : increment retrieved value; (result is 1)
- t_2 : decrement retrieved value; (result is -1)
- t_1 : store result in c ; (c is now 1)
- t_2 : store result in c ; (c is now -1)

This example shows that the race condition has caused the result from thread t_1 to be lost as it has been overwritten by the result from thread t_2 .

Racing for peripheral access

A race condition may also occur when two threads attempt to access the same peripheral, such as a printer spooler directory, at the same time.



In this example, the two processes *A* and *B* attempt to access the printer spooler directory at the same time:

- the next available printer job slot is 7;
- process *A* and *B* access printer job slot 7 simultaneously;
- process *A* reads printer job slot 7 and a timer interrupt occurs that causes a context switch to process *B* before process *A* has opportunity to store any data;
- process *B* reads printer job slot 7 and stores its job data and increments the values; and
- another timer interrupt occurs that causes a context switch to process *A* that then stores its job at printer job slot 7.

This example shows that the race condition has caused the job data stored in the printer spooler directory by process *B* to be lost as it has been overwritten by the job data from process *A*.

4.2 Inter-process synchronisation

Definition

Inter-process synchronisation involves techniques that are designed to prevent race conditions and allows threads/processes to share resources.

Behaviour of threads

Threads in a computer system may behave in two possible ways:

- competing – two or more processes compete for the same computing resource, for example access to a particular memory cell; or
- cooperating – two or more processes may need to communicate with one another, thus causing information to be passed from one to the other.

Inter-process synchronisation is required to manage both threads that are competing and threads that are cooperating.

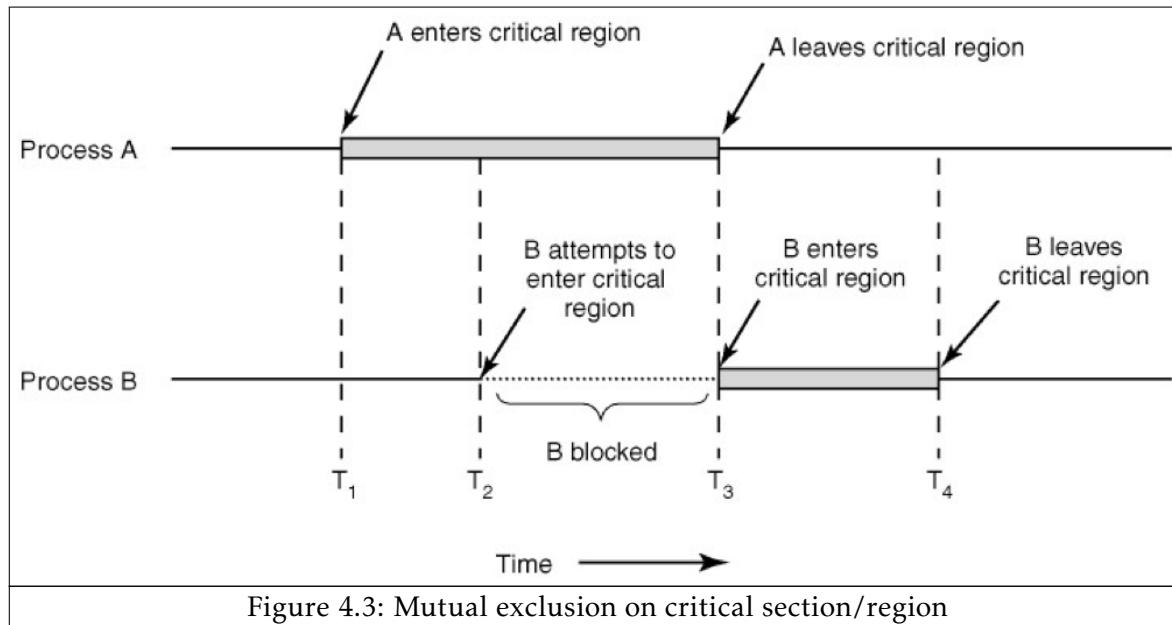
An operating system (OS) itself contains both threads that are competing and threads that are cooperating.

How mutual exclusion works

The solution to preventing race conditions is by implementing mutual exclusion on any given critical section/region.

The **critical section/region** is code in a process that involves sensitive operations on a shared resource.

Mutual exclusion is the requirement that one thread of execution never enters its critical section/region at the same time that another concurrent thread of execution enters its own critical section/region.



When a thread/process enters its critical section/region, no other thread/process may also enter its critical section/region.

This is demonstrated in the diagram above:

- at time interval T_1 , process A enters its critical section/region;
- at time interval T_2 , process B attempts to enter its critical section/region;
- process B is blocked from entering its critical section/region until process A leaves its critical section/region;
- at time interval T_3 , process A leaves its critical section/region and therefore process B is allowed to enter its critical section/region; and
- at time interval T_4 , process B leaves its critical section/region.

This shown that mutual exclusion is enforced as no two threads/processes are simultaneously inside their critical sections/regions.

For mutual exclusion to be effective:

- no assumptions may be made about the speeds or the number of threads/processes;
- no threads/processes running outside its critical section/region may block other threads/processes, such that a thread/process that is not in its critical section/region cannot prevent other threads from entering their critical section/region; and
- no thread/process should have to wait forever to enter its critical section/region.

Mutual exclusion is a major design issue in operating systems (OSs) as consideration must be taken in order to prevent race conditions while maintaining parallelism and efficiency.

How mutual exclusion is implemented

Mutual exclusion is implemented using semaphores.

A **semaphore** is a system tool used for the design of correct synchronisation protocols. This was introduced by Edsger Dijkstra in the 1962/1963. Semaphores are implemented using a variable or abstract data type and are used to control thread/process access to a resource. They are typically integer values that accept only non-negative values.



The diagram shows that semaphores allow the CPU to context switch between threads/processes when one becomes blocked.

It is convenient to write entry and exit protocols using a single atomic statement. This statement is atomic and therefore is indivisible, meaning that the statement cannot be interrupted.

As mentioned before, a semaphore, denoted by S , is an integer that takes only non-negative values. Only two atomic (indivisible) statements are permitted, as shown below.

Statement	Statement Implementation	Usage
$wait(s)$	<pre>wait(s) { if (S > 0) { S--; } }</pre>	If a thread/process is in its non-critical section/region and wishes to enter its critical section/region, this statement will be performed. This means that the thread/process will be blocked until $S = 0$ evaluates to <i>True</i> .
$signal(s)$	<pre>signal(s) { S++; }</pre>	If a thread/process is in its critical section/region, this statement will be performed. This helps to achieve mutual exclusion as it prevents $S = 0$ from evaluating to <i>True</i> until the thread/process has left its critical section/region.

This is a good solution as there is no possibility for a race condition as these statements will always be enforced due to the fact that they are atomic (indivisible) statements and cannot

be interrupted.

4.3 The producer-consumer problem

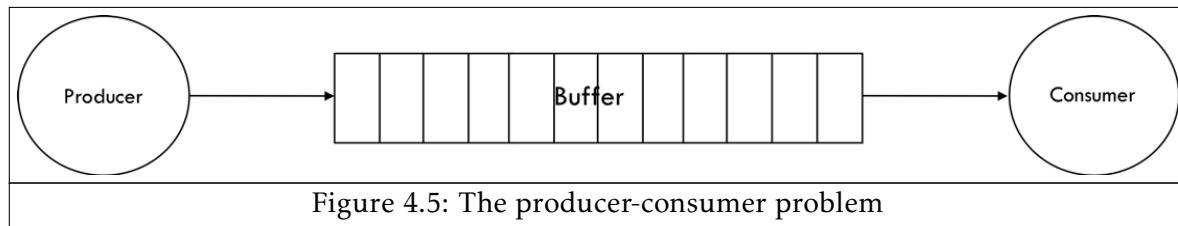
Problem description

The producer-consumer problem is a classical inter-process communication problem in which:

- a producer repeatedly produces items and places them in to a buffer; and
- a consumer consumes the items one-by-one by taking them from the buffer.

This problem has the following requirements:

- the buffer must be assumed to be first in, first out (FIFO);
- the producer may produce a new item only at a time when the buffer is not full;
- the consumer may consume an item only at a time when the buffer is not empty; and
- the process terminates when all items produced are eventually consumed.



The problem arises when attempting to devise a method that is able to:

- put the producer to “sleep” when the buffer is full to prevent further items being produced when there is no space in the buffer; and
- “wake” the consumer when the buffer is not empty as there is possibility to consumer when the buffer is not empty.

Possible solution

This problem could be solved by keeping track of the number of items in the buffer.

This could be achieved by implementing loops in the producer class and consumer class.

```

LOOP
{
    Produce item i          //produce item
    if ( itemCount == N )   //end of buffer
    {
        sleep(producer);
    }

    Put item i;            // place item in to buffer
    itemCount++;           // increment buffer count
    if ( itemCount == 1 )   // buffer nearly empty
    {
        wakeup(consumer);
    }
}

```

Producer class

```

LOOP
{
    if ( itemCount == 0 )   // buffer empty
    {
        sleep(consumer);
    }

    Remove item j;         // remove item from buffer
    itemCount--;           // decrement buffer count
    if ( itemCount == N-1 ) // buffer has space
    {
        wakeup(producer);
    }
    Consume item j;        // consume item
}

```

Consumer class

The loop in the producer class would be running as one thread and the loop in the consumer thread would be running as another thread. These two threads would be running in parallel. As a result, if the threads in the solution are interleaved, a race condition may occur, which in turn, may cause a deadlock.

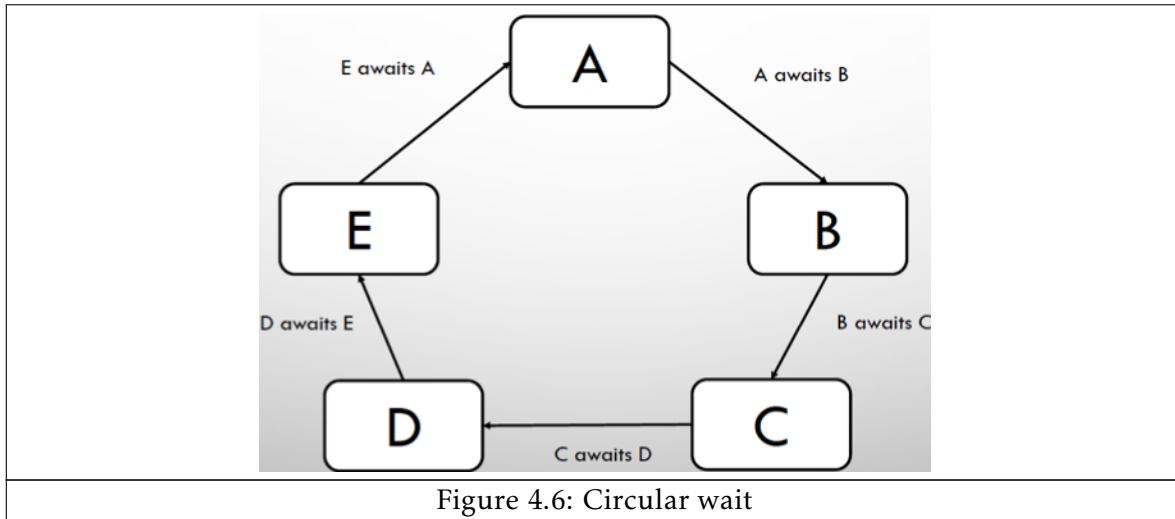
Deadlocks

A **deadlock** occurs when two or more threads wait for each other to finish.

Four conditions must be hold simultaneously in order for a deadlock to occur:

- mutual exclusion – a resource can be assigned to, at most, one process at a time;

- hold and wait – processes holding resources are permitted to request and wait for additional resources;
- no pre-emption – resources previously locked cannot be forcefully unlocked by another process, instead they must be released by the holding process; and
- circular wait – there must be a chain of processes, such that each member of the chain is waiting for a resource held by the next member of the chain, as shown in the diagram below.



A deadlock may occur in the possible solution described previously.

Consumer reads *itemCount* = 0 and it evaluates to *True*, and therefore *sleep(consumer)* needs to be called.

↓

Just before *sleep(consumer)* is called, the consumer is interrupted by a timer interrupt and the producer is resumed.

↓

The producer places an item in to the buffer, such that *itemCount* = 1.

↓

The producer tries to perform *wakeup(consumer)* however, the consumer is already in “wakeup” mode. As a result, the call to *wakeup(consumer)* is missed.

↓

When the consumer resumes, it will call *sleep(consumer)* and get trapped in “sleep” mode.

↓

The producer will continue placing items in the buffer and call *sleep(producer)* when the buffer is full.

↓

There is now a deadlock as both threads are waiting for a *wakeup* call from each other.

Figure 4.7: Possible deadlock

This possible deadlock shows that another solution is required to effectively solve the producer-consumer problem.

Solving the problem using semaphores

It is assumed that

```
ItemsReady = 0
SpacesLeft = N //size of buffer
```

```
LOOP
{
    Produce item i      // produce item
    Wait(SpacesLeft)   // decrement semaphore

    Put item i;        // place item in to buffer
    Signal(ItemsReady) // increment
}
```

Producer class

```
LOOP
{
    Wait(ItemsReady)   // decrement semaphore

    Get item j;        // remove item from buffer
    Signal(SpacesLeft) // increment semaphore
    Consume item i;   // consume item
}
```

Consumer class

If this solution uses semaphores correctly, then

$$N = \text{SpacesLeft} + \text{ItemsReady}$$

as the producer will always be placing items in to the buffer when there are spaces available in the buffer.

However, this solution does not consider situations in which there are multiple producers and/or multiple consumers.

The multiple producer-consumer problem

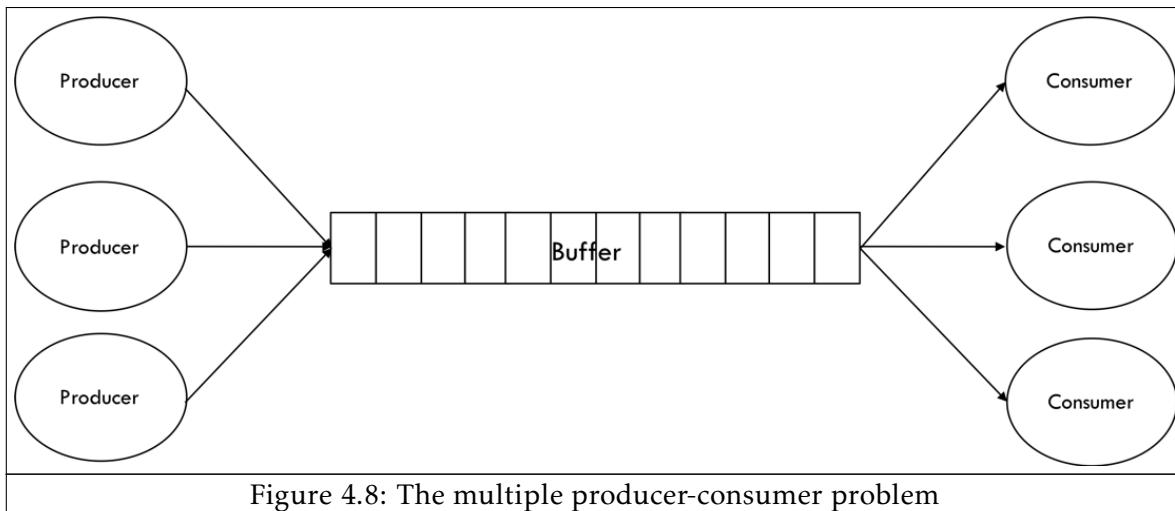
Problem description

The multiple producer-consumer problem is a classical inter-process communication problem in which:

- multiple producer repeatedly produces items and places them in to a buffer; and
- multiple consumer consumes the items one-by-one by taking them from the buffer.

As with the previous producer-consumer problem, this problem has the following requirements:

- the buffer must be assumed to be first in, first out (FIFO);
- the producers may produce a new item only at a time when the buffer is not full;
- the consumers may consume an item only at a time when the buffer is not empty; and
- the process terminates when all items produced are eventually consumed.



The problem arises when attempting to devise a method that is able to manage:

- two producers placing items in to the same slot in the buffer; and
- two consumers removing items from the same slot in the buffer.

This is similar to the problem discussed in the printer spooler example (page n).

A race condition may also occur when producers attempt to access a variable at the same time.

To demonstrate the race condition, it is necessary to consider the following possible interleaving of the threads/processes:

- two producers access the *SpacesLeft* variable at the same time, which corresponds to decrementing the semaphore;

- both producers get the same next empty slot in the buffer at the same time; and
- both producers write in to the same slot.

This example shows that the race condition has caused the data stored in the buffer slot by the first producer to be lost as it has been overwritten by the data stored in the buffer slot by the second producer.

In order to ensure mutual exclusion when multiple users are involved, an additional semaphore must be introduced.

Mutex

A **mutex** (or **binary semaphore**) is a semaphore with ownership that can only be released by its owner and is initially set to 1.

Problem solution

It is now possible to construct a solution, using a mutex (or binary semaphore), that will ensure mutual exclusion even when there are multiple producers and/or multiple consumers.

It is assumed that

```
ItemsReady = 0
SpacesLeft = N //size of buffer
```

```
LOOP
{
    Produce item i      // produce item
    Wait(SpacesLeft)   // decrement semaphore
    Wait(BusyBuffer)   // mutex

    Put item i;        // place item in to buffer
    Signal(BusyBuffer) // release mutex
    Signal(ItemsReady) // increment
}
```

Producer class

```
LOOP
{
    Wait(ItemsReady)   // decrement semaphore
    Wait(BusyBuffer)   // mutex

    Get item j;        // remove item from buffer
    Signal(SpacesLeft) // increment semaphore
    Signal(BusyBuffer) // release mutex
```

```
    Consume item i;      // consume item  
}
```

Consumer class

The mutex *BusyBuffer* has ownership and therefore can only be incremented/decremented by the same thread/process.

The order in which semaphores are incremented and decremented is essential. This can be demonstrated by inspecting the effect of switching around two statements in the Consumer class:

```
Wait(ItemsReady) //decrement semaphore  
Wait(BusyBuffer) //mutex  
...  
Wait(BusyBuffer) //mutex  
Wait(ItemsReady) //decrement semaphore
```

This switching would cause ...

4.4 The dining philosophers problem

Problem description

The dining philosophers problem is a classical inter-process communication problem in which:

- five philosophers are seated around a circular table eating and thinking; and
- each philosopher has a plate of spaghetti that they can eat with forks.

This problem has the following requirements:

- the life of a philosopher consists of only alternating periods of eating and thinking;
- between each pair of plates is one fork;
- each philosopher needs two forks to eat the spaghetti; and
- no two philosophers may hold the same fork simultaneously.

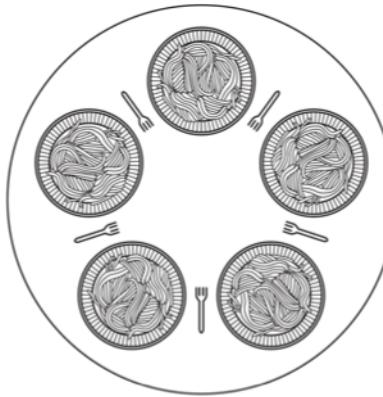


Figure 4.9: Layout of table

The problem arises when attempting to devise a method that is able to:

- allow each philosopher to have alternating periods of eating and thinking; and
- not result in a deadlock.

It could be said that the problem requirement for each philosopher to need two forks to eat the spaghetti is somewhat artificial. As a result, we can substitute the spaghetti for rice and substitute chopsticks for forks.

The problem now has the following updated requirements:

- the life of a philosopher consists of only alternating periods of eating and thinking;
- between each pair of plates is one chopstick;
- each philosopher needs two chopsticks to eat the rice; and
- no two philosophers may hold the same chopstick simultaneously.

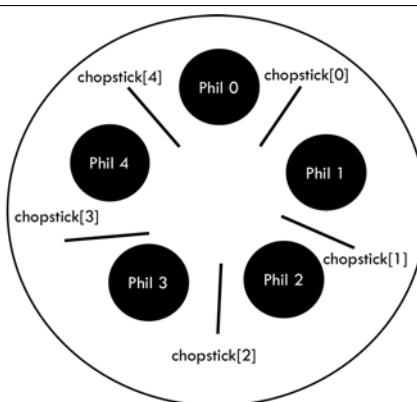


Figure 4.10: Updated layout of table

Problem solutions

Problem solution 1

This problem can be solved using semaphores, using the following assumptions:

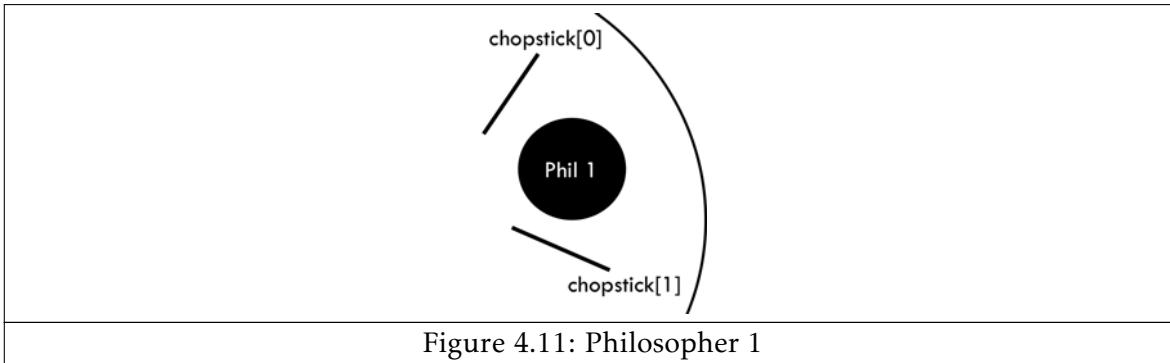
- each philosopher is a different thread with a unique ID;
- one semaphore is implemented per chopstick; and
- chopsticks are identified by using the unique ID of a philosopher.

As chopsticks are identified by using the unique ID of a philosopher, it could be that:

- the chopstick to the left of a given philosopher is $\text{chopstick}[i]$; and
- the chopstick to the right of a given philosopher is $\text{chopstick}[(i - 1) + N] \% N$.

where i is the ID of the philosopher and N is the number of philosophers.

The identification of the chopsticks works as demonstrated by using Philosopher 1 as an example.



The diagram shows that:

- the chopstick to the left of Philosopher 1 is $\text{chopstick}[1]$; and
- the chopstick to the right of Philosopher 1 is $\text{chopstick}[0]$.

Using this example, it is possible to validate the chopstick identification formulas discussed before.

Given that $i = 1$ for Philosopher 1 and that $N = 5$ as there are five philosophers in total,

```

left = chopstick[i]
left = chopstick[1]
right = chopstick[(i-1) + N] % N
right = chopstick[(1-1) + 5] % 5
right = chopstick[(0 + 5) % 5]
right = chopstick[5 % 5]
right = chopstick[0]

```

This shows that the chopstick identification formulas work for Philosopher 1.

```
LOOP
{
    think();
    wait(chopstick[left]);      // take left chopstick
    wait(chopstick[right]);     // take right chopstick
    eat();
    signal(chopstick[left]);   // release left chopstick
    signal(chopstick[right]);  // release right chopstick
}
```

Producer class

However, this solution has the possibility to cause a race condition.

This can be demonstrated in the situation where all five philosophers wish to eat at the same time and therefore all take their left chopsticks:

- Philosopher 0 takes *chopstick[left]*;
- Philosopher 1 takes *chopstick[left]*;
- Philosopher 2 takes *chopstick[left]*;
- Philosopher 3 takes *chopstick[left]*; and
- Philosopher 4 takes *chopstick[left]*.

This causes a race condition as:

- each philosopher will now be waiting to take *chopstick[right]*;
- no philosopher can take *chopstick[right]* as this chopstick is the subsequent philosopher's *chopstick[left]* and this has already been taken; and
- no philosopher can perform their *eat* function and therefore no chopsticks will be released as the *signal* functions are only performed after the *eat* function has completed.

This shows that, if the threads in the solution are interleaved, a race condition may occur. In this case, a circular wait is caused and therefore there is a deadlock.

Possible solution 2

As shown in the multiple producer-consumer problem, a mutex (or binary semaphore) can be introduced to prevent this deadlock.

```
LOOP
{
    think();
    wait(busy);                // mutex
    wait(chopstick[left]);     // take left chopstick
    wait(chopstick[right]);    // take right chopstick
```

```

eat();
signal(chopstick[left]);    // release left chopstick
signal(chopstick[right]);  // release right chopstick
signal(busy);              // release mutex
}

```

Producer class

This solves the deadlock as the mutex (or binary semaphore) signals the critical section/region and prevents other philosophers from attempting to take a chopstick that is currently in use by another philosopher.

Although this solution is deadlock-free, it has a performance bug. The mutex means that only one philosopher can be eating at any instant and, with five chopsticks available, it should be possible to allow two philosophers to eat at the same time. As a result, there are more efficient solutions to this problem that achieve maximum parallelism.

Final revision solution

This solution is deadlock-free and allows the maximum parallelism for an arbitrary number of philosophers.

```

#define N      5           // number of philosophers
#define LEFT   (i + N - 1) % N // ID of i's left neighbour
#define RIGHT  (i + 1) % N   // ID of i's right neighbour
#define THINKING 0          // philosopher is thinking
#define HUNGRY   1          // philosopher is trying to acquire
                        // chopsticks
#define EATING   2          // philosopher is eating

typedef int semaphore           // semaphores are a special kind of
                                // integer
int state[N];                  // array to keep track of philosopher's
                                // state
semaphore mutex = 1;           // mutual exclusion for critical regions
semaphore s[N];                // one semaphore per philosopher

// i: philosopher unique ID, from 0 to N-1
void philosopher (int i)
{
    while (TRUE)               // repeat indefinitely
    {
        think();                // philosopher is thinking
        take_chopsticks(i);     // acquire two chopsticks or be blocked
        eat();                   // philosopher is eating
        put_chopsticks(i);      // place both chopsticks back on the table
    }
}

```

```

// i: philosopher unique ID, from 0 to N-1
void take_chopsticks(int i)
{
    wait(&mutex);                      // enter critical section/region
    state[i] = HUNGRY;                  // record fact that philosopher i is
    hungry
    test(i);                          // attempt to acquire two chopsticks
    signal(&mutex);                  // exit critical section/region
    wait(&s[i]);                     // block if chopsticks were not acquired
}

// i: philosopher unique ID, from 0 to N-1
void put_chopsticks(int i)
{
    wait(&mutex);                      // enter critical section/region
    state[i] = THINKING;               // record fact that philosopher i is
    // thinking
    test(LEFT);                       // check if left neighbour can now eat
    test(RIGHT);                      // check if right neighbour can now eat
    signal(&mutex);                  // exit critical section/region
}

// i: philosopher unique ID, from 0 to N-1
void test(int i)
{
    if (state[i] == HUNGRY && state[LEFT] != EATING
        && state[RIGHT] != EATING)
    {
        state[i] = EATING;           // record fact that philosopher i is hungry
        signal(&s[i]);              // exit critical section/region
    }
}

```

Producer class

This solution uses the array state to keep track of whether a philosopher is:

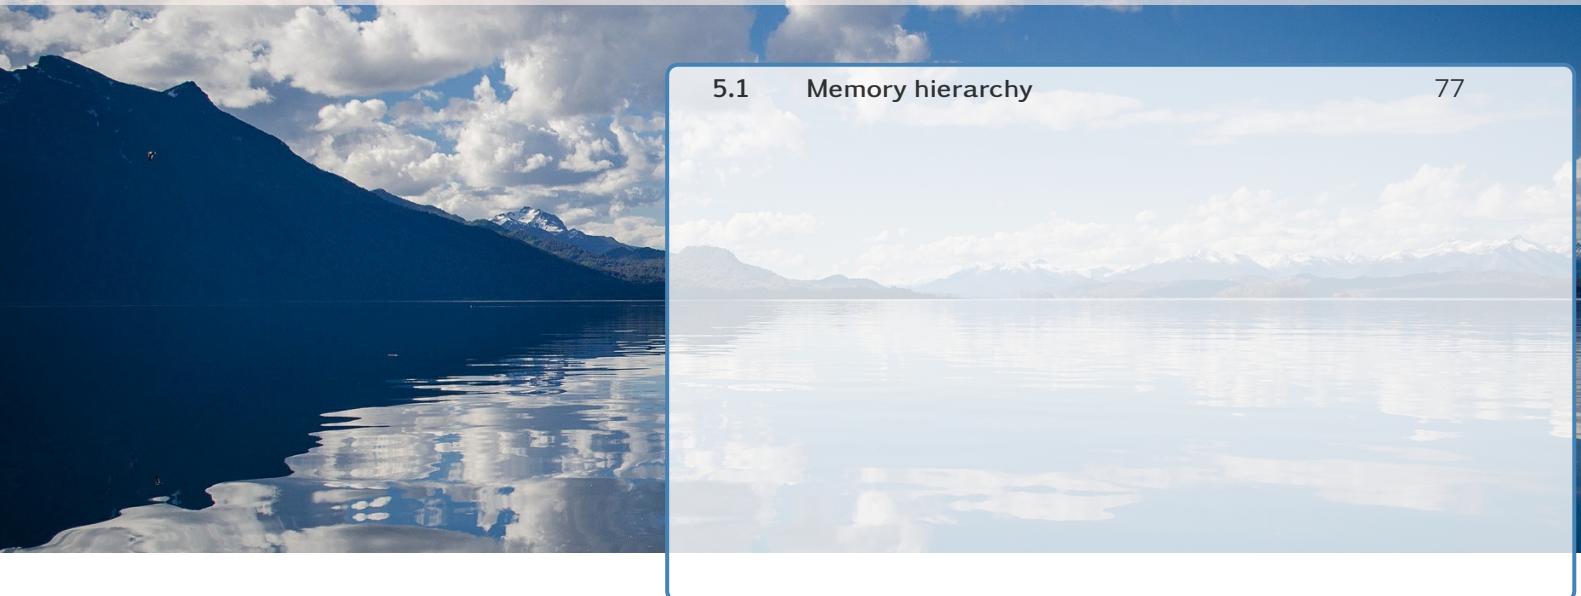
- eating;
- thinking; or
- hungry (trying to acquire chopsticks).

This represents an array of semaphores and each philosopher has its own *state* array. This enables hungry philosophers to be blocked if the required chopsticks are currently busy.

A philosopher may move in to eating state only if neither neighbouring philosopher is eating. For example, Philosopher 1 cannot enter eating state if:

- Philosopher 0 is currently in eating state; or
- Philosopher 2 is currently in eating state.

5. Memory Management



5.1 Memory hierarchy

77

5.1 Memory hierarchy

Definition

The **memory hierarchy** separates computer storage into a hierarchy based on response time. A computer system is usually composed by a layered memory system

Diagram

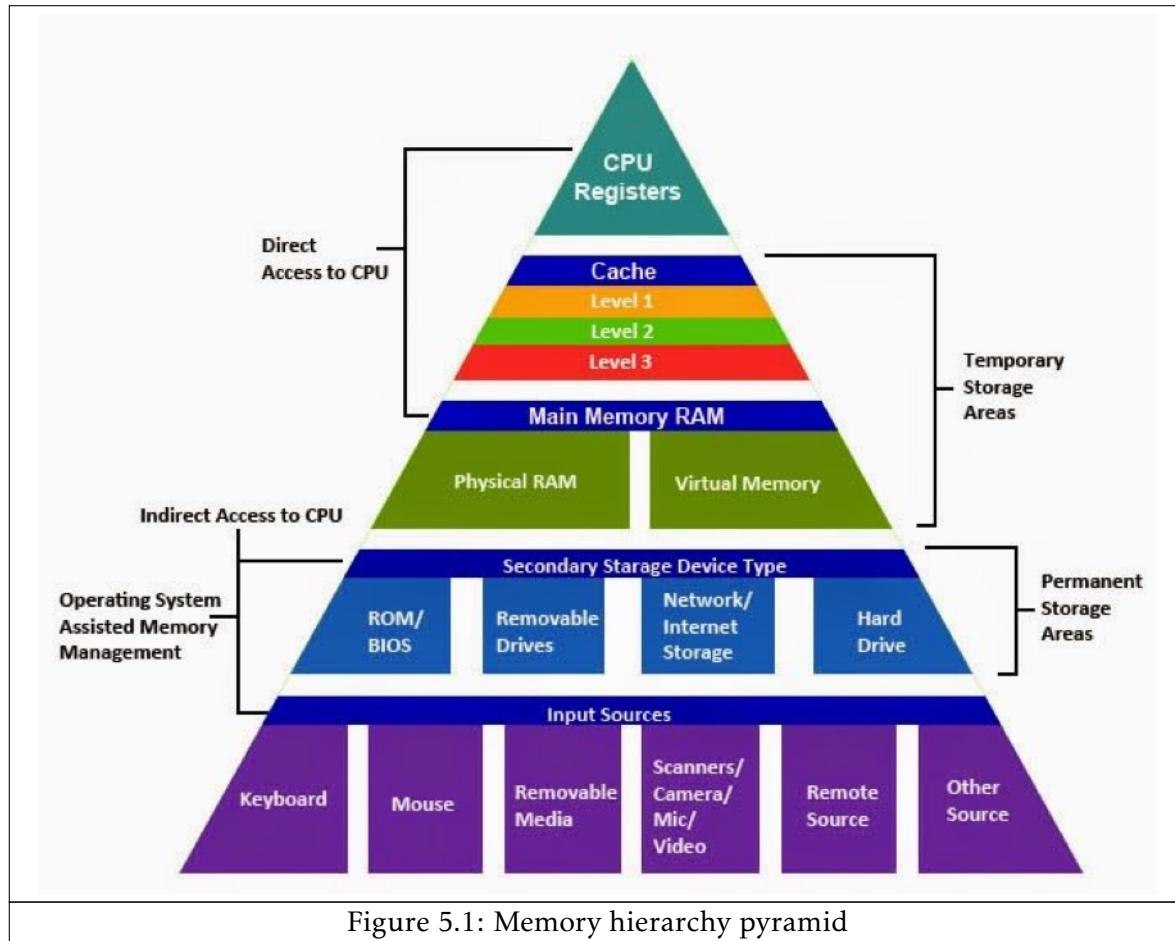


Figure 5.1: Memory hierarchy pyramid

Higher layers correspond to faster devices that:

- have lower capacity;
- require more power and generate more heat; and
- are more expensive.

Lower layers correspond to slower devices that:

- have higher capacity;
- require less power and generate less heat; and
- are cheaper.

The temporary storage areas can be described as volatile; the data does not persist after power-down.

The permanent storage areas can be described as non-volatile; the data persists after power-down.

